

Dynamic Load Balancing and Persistent Chat

CSD Course — Lab Assignment 6 (updated task)

Student Name	Rahul Raj
Roll Number	12341680
Date	10 September 2026
Load balancer URL	<code>http://10.1.75.53:3269</code>
Required routes	<code>POST /message · GET /feed</code>
Source code	<code>https://github.com/Rahul5977/csd-lab6-dynamic-lb-persistent-chat</code>

1. Submission Details

The only endpoint clients need is the load balancer. Backends are never addressed directly.

Load balancer URL	http://10.1.75.53:3269
Submit a message	POST http://10.1.75.53:3269/message with client-name and msg
Read the feed	GET http://10.1.75.53:3269/feed
Chat application	http://10.1.75.53:3269/ — the full secure group chat (register, log in, end-to-end encrypted rooms)
Balancer dashboard	http://10.1.75.53:3269/lb/ · JSON at /lb/stats , /lb/events
Repository	https://github.com/Rahul5977/csd-lab6-dynamic-lb-persistent-chat

```
curl -X POST http://10.1.75.53:3269/message \
  -H 'Content-Type: application/json' \
  -d '{"client-name": "rahul", "msg": "hello"}'
# {"ok":true,"duplicate":false,"id":"635e3754-...", "seq":1,"client-
name":"rahul","backend":"sys4"}

curl http://10.1.75.53:3269/feed
# {"ok":true,"room":"public","backend":"sys2","count":20602,"returned":200,"truncated":true,...}
```

/message also accepts a form-encoded body or query parameters, and GET as well as POST, because the encoding used by the evaluation load generator is not specified. Both routes are additions to the application: **nothing was removed or simplified**. The authenticated chat, the script password hashing, the sessions and the end-to-end-encrypted locked rooms of the previous assignment all still work through the same URL and are still covered by the test suites (§9).

2. Assigned Systems and Ports

Role	Alias	SSH	Container IP	Service port(s)	External port
Load balancer	stu78_sys1	ssh -p 2269 student@10.1.75.53	172.17.0.70	LB 3000	3269 → the only URL clients use
Backend A	stu78_sys2	ssh -p 2270 student@10.1.75.53	172.17.0.71	app 3000	3270 (direct, kept for the previous assignment)
Backend B + database	stu78_sys3	ssh -p 2271 student@10.1.75.53	172.17.0.72	app 3000 · SQLite database service 5270	3271 (direct, previous assignment)
Backend C	stu78_sys4	ssh -p 2272 student@10.1.75.53	172.17.0.73	app 3001	— (port 3000 / external 3272 belongs to my course project)

The four systems are Ubuntu 24.04 containers on one shared lab host, each limited by cgroup to **exactly one CPU core** and 512 MB. The lab NAT forwards only the first external port of each system to container port 3000. Clients talk only to 10.1.75.53:3269; the load balancer reaches the backends over the private bridge network, so a backend needs no public port at all — which is why sys4's backend can live on 3001. The load generator runs on my own machine.

The database service started out next to the load balancer on sys1. Section 13 shows the measurement that moved it to sys3: the two processes were fighting over sys1's single CPU and the container was being CFS-throttled in half of all scheduling periods. Moving it was worth **+41 % throughput**.

3. Objective


The task was set in two parts. The first asked for a dynamically load-balanced, persistent, duplicate-proof group chat; the update fixed the public API, demanded that switching be driven by a *threshold* on backend load, and asked for that threshold to be optimised and for the report to plot the utilisation of all four systems.

Original. Extend the previous secure group-chat application and deploy its backend on the allotted systems behind a load balancer. The balancer must select a backend **dynamically** using response time and/or current system load, monitor backend health and stop sending to unavailable backends, and **automatically detect backends started while the application is running**. All backends must share **persistent** chat data, every message must have a **unique id**, and duplicate insertion must be prevented on retries and reconnections.

Updated. Host the application through the load balancer so clients use only its URL, and submit that URL. Switch backends **when the load of the current backend exceeds a defined threshold** — fixed round robin is not acceptable — and **determine the optimal threshold**. Expose exactly `/message` (taking `client-name` and `msg`) and `/feed` (retrieving messages). Write an own load generator supporting a **variable number of users, random message lengths and random intervals**, and report plots including response time and **the system utilisation of all four systems**. The submitted application must remain the previous secure, persistent group chat and **must not be simplified to chase leaderboard rank**.

Everything above is implemented, deployed on the allotted systems, measured and demonstrated in this report. Section by section: the required routes are §5, the threshold rule is §6.1, the optimisation of the threshold is §10, the response-time and utilisation plots are §11, and §14 states exactly what was added and what was left untouched.

4. Architecture



Request lifecycle. A request arrives at the balancer, which picks a backend with the threshold rule (§6.1), adds `X-Forwarded-For` / `X-Real-IP` and proxies it over a pooled keep-alive connection. The backend performs the operation against the database service and answers with `X-Backend-Id` (which instance served it); the balancer adds `X-LB-Active-Backends` (how many backends were routable at that moment), and the

load generator records both on every request. A new message is appended once to SQLite; the database broadcasts it on a WebSocket firehose and every backend pushes it to its own connected browsers, so three stateless backends behave as one application. Browser WebSockets are pinned to a backend by IP-hash and tunnelled byte-for-byte.

What is inherited from the previous assignments unchanged: bcrypt password hashing with constant-time compare and login rate limiting; end-to-end encrypted "locked" rooms (PBKDF2-HMAC-SHA-256 with 250 000 iterations → AES-256-GCM, fresh IV per message, AAD binding room and key epoch, key derived in the browser — the server stores only ciphertext); the WebSocket delivery path; the single-page UI.

5. Required API Routes — `/message` and `/feed`

The update fixes two route names on the balancer. They are **new entry points into the existing application**, not a new application: same database, same room, same message-id duplicate guard, same live delivery. A message posted to `/message` appears immediately in an open browser chat, and a message typed in the browser appears in `/feed`.

5.1 POST `/message` — "client-name" and "msg"

The evaluation generator's request encoding is unspecified, so the route is deliberately permissive and answers the same JSON either way:

Accepted	Detail
Method	POST (also GET, for a generator that only builds URLs)
Body	JSON, <code>application/x-www-form-urlencoded</code> , a raw text body, or query-string parameters
Client name	<code>client-name</code> , and the aliases <code>client_name</code> , <code>clientName</code> , <code>name</code> , <code>user</code> , <code>username</code> , <code>from</code> , <code>sender</code> ; sanitised to 48 characters, empty becomes <code>anonymous</code>
Message	<code>msg</code> , and the aliases <code>message</code> , <code>text</code> , <code>body</code> , <code>content</code> ; trimmed, capped at 2 000 characters
Message id	optional — <code>id</code> , <code>msg-id</code> , <code>message-id</code> , <code>uuid</code> , or the <code>Idempotency-Key</code> header. If none is supplied the backend mints a UUID v4. An id that does not fit the id grammar is hashed into one rather than rejected, so it is still one id per client message

The reply names the id, the per-room sequence number, whether it was a duplicate, and which backend served it:

```
{"ok":true,"duplicate":false,"id":"635e3754-cedb-42fc-9b71-2ca04479e208",
"seq":1,"client-name":"rahul","client_id":false,"room":"public","backend":"sys4"}
```

A repeat of the same id is answered 200 `{"ok":true,"duplicate":true,...}` with header `X-Duplicate: 1` and the *original* sequence number — an idempotent success, never an error the client would retry again. The `dedup` field names which of the three layers caught it (§8.3).

5.2 GET `/feed` — the messages

`/feed` returns the shared room, newest-last, with the sender, the id, the sequence number and the serving backend on every row.

```
{"ok":true,"room":"public","backend":"sys2","count":20602,"returned":200,
"truncated":true,"limit":200,"messages":[{"id":"...", "seq":20403, "from":"client-0007",
"ts":1789021..., "kind":"plain", "via":"sys3", "text":"...", ...}]}
```

Why the default is a window and not the entire history. The room grows without bound under a load generator, and the cost of returning all of it is linear in its size: measured on the live system, at 1 412 messages the full feed was already 383 KB and 190 ms, and it only gets worse. A feed endpoint that returns

everything therefore ends up measuring the size of the history rather than the system. The default answer is the newest **200** messages — what a chat client actually renders — and every answer states the true total in `count`, whether it was `truncated`, and the `limit` used. `?limit=N` and `?limit=all` return more, up to the **complete history**, so no message is unreachable:

```
curl 'http://10.1.75.53:3269/feed?limit=all' | jq '.count, .returned'
```

Each backend caches the default body and rebuilds it at most every 100 ms — invalidated the instant *any* backend appends, because the database firehose tells all of them — so a read-heavy generator cannot make the database re-serialise the same window thousands of times a second. The database applies the same rate cap to its own serialisation, and keeps the room's row count incrementally instead of running `COUNT(*)` per request.

6. Load Balancer Design (sys1)

`lb/loadbalancer.py` — pure Python 3 standard library, no pip, no root, no external process. It is the previous assignment's balancer extended in four areas: performance-based **threshold** selection, three-state health monitoring, live membership, and observability.

6.1 Performance-based selection: the threshold rule

"The load balancer must dynamically select backends based on system performance/load. When the load of the current backend exceeds a defined threshold, traffic must switch to another suitable backend. Fixed round-robin alone is not acceptable."

Each backend gets a single **load index** between 0 (idle) and 1 (saturated) — the worst of three signals, so no single blind spot can hide a busy backend:

```
load(b) = max(  cpu(b)                the container's own cgroup CPU utilisation, read
                inflight(b) / inflight_cap  from the backend's /health (at most one probe old)
-             requests this balancer is holding on it right now
                ewma_rt(b) / rt_cap_ms )    measured locally, so it reacts instantly
budget                                     EWMA response time against the response-time
```

with `inflight_cap = 24` and `rt_cap_ms = 250`; a DEGRADED backend is forced to at least 1.0. The selection rule is then literally the sentence in the task:

Condition	Action
<code>load(current) < T</code>	keep sending to the current backend
<code>load(current) ≥ T</code>	switch to a backend still under <code>T</code>
no backend under <code>T</code>	send to the least-loaded one — degrade, never refuse

The replacement is chosen by **power-of-two-choices** among the candidates under `T`: two are sampled at random and the better one wins. Picking the single global minimum instead would make every concurrent request that crosses the threshold jump to the *same* replacement, which would then cross the threshold itself — classic herding. Sampling two spreads the switch without any coordination.

Two details make the rule react fast enough to matter. The **in-flight term is measured by the balancer itself**, so a backend that starts queueing crosses the threshold within one request rather than waiting up to a health interval for a CPU number. And a **never-sampled backend is selected immediately**, so a backend started while the application is running receives traffic on its first eligible request instead of waiting for a probe history.

Deliberately, this is *not* round robin: below the threshold an idle cluster concentrates work on one warm backend, which keeps its connection pool and caches hot, and the *load itself* — not a counter — decides when to spread out. `switch_threshold` is tunable at runtime through `POST /lb/config`, which is what the sweep in §10 drives.

6.2 The adaptive score — kept as the comparison baseline

The previous algorithm is still available and is the baseline in §11.4:

```
score(b) = EWMA_rt(b) × (1 + in_flight(b)) × (1 + load_weight × cpu_load(b)) / weight(b)
```

also with power-of-two-choices, plus `least_response_time` and the four classic algorithms (`round_robin`, `least_connections`, `weighted_round_robin`, `ip_hash`). Nothing was removed; the threshold rule was added and made the default.

6.3 Health monitoring

A background thread probes every backend's `/health` every 3 s (6 s timeout) and reads its load JSON. States instead of a boolean:

State	Meaning	Routable?
UP	answering in time	yes
DEGRADED	answering, but probe slower than 2 s or reported CPU ≥ 95 %	yes, but forced over the threshold
DOWN	connection refused / timeouts for 2 consecutive checks; back to UP after 2 successes	no
DRAINING	asked to leave (<code>/lb/deregister</code>): finishes in-flight, gets nothing new	no

Passive checks eject a backend immediately when a live proxy attempt cannot connect; the request is then retried on another backend, and only if nothing was sent upstream yet — so the balancer can never duplicate a POST, and even if it did the message id would make it harmless. **Fail-open:** the last routable backend is never ejected. "Slow ≠ dead" fixes the ejection churn diagnosed in the previous assignment, where a saturated one-core backend answered its health check late, was ejected, and the survivors then collapsed in turn.

6.4 Dynamic membership — detecting backends started while running

Three mechanisms, any one of which is sufficient:

1. **Self-registration (push).** Each backend POSTs `/lb/register {id, host, port, weight}` with a shared token at boot and every 5 s as a heartbeat. An unknown backend is admitted on the spot. On SIGTERM it POSTs `/lb/deregister` and drains — graceful scale-down.
2. **Candidate scan (pull).** The configuration lists candidate slots (sys3:3000, sys4:3001). Every 5 s the balancer probes the ones it does not know and admits any that answers `/health` with the expected application version. During development this correctly *refused* the previous assignment's backend still running on sys3, because it reports a different version.
3. **Config watch.** The configuration file is re-read whenever its modification time changes (also `SIGHUP` and `POST /lb/reload`); `POST /lb/config` changes the algorithm or the threshold at runtime.

Dynamically discovered backends that stay DOWN for 15 minutes are pruned; statically configured ones are kept as DOWN.

6.5 Observability

`/lb/stats` reports per-backend state, source, score, **load index**, EWMA, probe latency, load, in-flight, counts and percentiles, plus `active_backends`, the current `switch_threshold`, which backend traffic is currently pinned to and how many times the threshold rule has switched. `/lb/events` is the timeline of added /

degraded / ejected / readmitted / draining / switch / config events; /lb/ is the auto-refreshing dashboard (§15); and the CSV access log carries the active-backend count on every line.

7. Load Balancer — Key Code

The complete balancer (lb/loadbalancer.py , ~830 lines) is in the repository on the cover page. The two parts that implement the task's requirement are the load index and the threshold rule:

```
1  def cpu_load(self):
2      """0..1 : how busy the backend's SYSTEM is. Prefers the container-wide
3      cgroup CPU utilisation (sys_cpu_pct - sees other tenants of the box),
4      falls back to the process CPU % and the load average per core."""
5      cpu = float(self.load.get("cpu_pct", 0) or 0) / 100.0
6      sys_cpu = self.load.get("sys_cpu_pct")
7      if sys_cpu is not None:
8          cpu = max(cpu, float(sys_cpu) / 100.0)
9      cores = max(1, int(self.load.get("cores", 1) or 1))
10     la = float(self.load.get("loadavg1", 0) or 0) / cores
11     return max(cpu, min(la, 4.0))
12
13  def load_index(self, cfg):
14      """0 = idle, 1 = saturated. The single number the threshold rule compares
15      against 'switch_threshold'. Three signals, worst one wins:
16      cpu        system/cgroup CPU from the backend's /health (<= health_interval old)
17      queue      in-flight requests THIS LB is holding on it (instantaneous)
18      latency    EWMA response time against the response-time budget
19      The queue term is what makes the rule react inside one health interval:
20      a backend that starts queueing crosses the threshold immediately."""
21      cpu = self.cpu_load()
22      queue = self.in_flight / max(1.0, float(cfg["inflight_cap"]))
23      lat = 0.0 if self.ewma_ms is None else self.ewma_ms / max(1.0, float(cfg["rt_cap_ms"]))
24      idx = max(cpu, queue, lat)
25      if self.state == DEGRADED:
26          idx = max(idx, 1.0) # degraded is by definition over threshold
27      return idx / max(1, self.weight)
28
29  def score(self, cfg):
30      """Lower is better. New backends (no sample) score 0 -> tried at once."""
31      if self.ewma_ms is None:
32          return 0.0
33      s = self.ewma_ms * (1 + self.in_flight)
34      if cfg["algorithm"] == "adaptive":
35          s *= 1 + cfg["load_weight"] * self.cpu_load()
36      if self.state == DEGRADED:
37          s *= 2.0
38      return s / max(1, self.weight)
```

```
1  def _pick_threshold(self, pool):
2      """THRESHOLD algorithm - the updated assignment's explicit requirement:
3      keep sending to the current backend while its load stays below a defined
4      threshold, and switch to another suitable backend as soon as it does not.
5
6      load_index(current) < T  -> keep using it
7      load_index(current) >= T -> switch to a backend still under T
8                               (power-of-two-choices among them, so
9                               concurrent requests do not all herd
10                              onto the same replacement)
11      no backend under T      -> everything is hot: send to the lowest
12                              score, i.e. degrade gracefully rather
13                              than refuse traffic
14
15      This is deliberately not round-robin: an idle cluster concentrates work
16      on one warm backend (better cache locality, fewer cold connections) and
17      the load itself, not a counter, decides when to spread out."""
18      cfg = self.cfg
19      T = float(cfg["switch_threshold"])
20      fresh = [b for b in pool if b.ewma_ms is None]
21      if fresh: # a newly discovered backend: measure it now
22          self.current_id = fresh[0].id
23          return fresh[0]
24      if len(pool) > 1 and random.random() * 100 < cfg["explore_pct"]:
25          return random.choice(pool) # keeps every backend's estimate fresh
26      cur = next((b for b in pool if b.id == self.current_id), None)
```

```

27     if cur is not None and cur.load_index(cfg) < T:
28         return cur # under threshold -> stay put
29     if cur is not None and cfg["min_dwell_ms"] and \
30         (time.time() - self.last_switch) * 1000 < cfg["min_dwell_ms"]:
31         return cur # hysteresis: too soon to move again
32     under = [b for b in pool if b.load_index(cfg) < T and b is not cur]
33     # Power-of-two-choices in BOTH branches. Picking the single global minimum
34     # would make every concurrent request that crosses the threshold jump to the
35     # same replacement, and that backend would then cross the threshold itself -
36     # sampling two candidates and keeping the better one spreads the switch.
37     candidates = under or [b for b in pool if b is not cur] or pool
38     a, b = random.choice(candidates), random.choice(candidates)
39     key = (lambda x: (x.score(cfg), x.in_flight)) if under \
40         else (lambda x: (x.load_index(cfg), x.score(cfg), x.in_flight))
41     chosen = min((a, b), key=key)
42     if chosen.id != self.current_id:
43         self.switches += 1
44         self.last_switch = time.time()
45         if self.switches % 25 == 1: # timeline for the report, not a log flood
46             self.event("switch", chosen.id,
47                 "threshold %.2f exceeded on %s" % (T, cur.id if cur else "-"))
48     self.current_id = chosen.id
49     return chosen

```

8. Database Persistence and Duplicate Prevention

8.1 Persistent shared storage

All backends are stateless; every user, session, room and message lives in **one SQLite database (WAL mode)** owned by `app/db_service.js` on `sys3` and reached over HTTP by every backend. The database file (`~/assignment6/data/chat.sqlite`) survives restarts of any backend, of the database service, and of the whole cluster (verified live in §12.3). The previous assignment's data (717 users, 5 rooms, 111 180 messages) was migrated in on first boot, so the old accounts still work, and the file was carried across intact when the service moved from `sys1` to `sys3` — WAL checkpointed first, no rows lost. SQLite was chosen because the lab boxes allow no root and no package installation; Node 22's built-in `node:sqlite` needs nothing but a user-local Node binary, and a real relational engine gives real constraints.

8.2 Unique message ids

Every message has a **UUID v4 id minted by the client when the message is composed** and reused on every retry (body field `id` or the `Idempotency-Key` header). Only the client knows that two sends are "the same message" — a server-minted id cannot recognise a retry — so this is the right place for the id. A client that sends none, which includes the evaluation load generator, gets one minted by the backend: still globally unique, but then a retry is by definition a new message.

8.3 How duplicates are prevented — three layers

Layer	Where	Mechanism
1 · fast path	each backend	LRU of the last 5 000 stored ids, also fed by the firehose, so a retry that lands on a <i>different</i> backend is usually caught here too → <code>200 {duplicate:true}</code> without touching the database
2 · authoritative	database	<code>messages.id</code> TEXT PRIMARY KEY; POST /messages runs <code>SELECT ... ; INSERT ... ON CONFLICT(id) DO NOTHING</code> inside one <code>BEGIN IMMEDIATE</code> transaction. Two backends sending the same id at the same instant are serialised; the loser gets the stored row back with <code>duplicate:true</code> and the rejection is written to <code>dedup_log</code>
3 · presentation	browser	renders each id once; a retried send that the server deduplicated never appears twice

A duplicate is answered with **HTTP 200**, `duplicate:true`, the original `seq` and `id`, and header `x-Duplicate: 1` — an idempotent success, never an error the client would retry again. Per-room ordering uses `UNIQUE(room, seq)` so sequences are gap-free. The append transaction:

```

1  const appendTx = (room, entry) => {
2    db.exec('BEGIN IMMEDIATE');
3    try {
4      const existing = q.getMsg.get(entry.id);
5      if (existing) {
6        q.logDup.run(entry.id, room, entry.via || null, Date.now());
7        db.exec('COMMIT');
8        return { duplicate: true, rec: rowToEntry(existing) };
9      }
10     const seq = q.nextSeq.get(room).seq;
11     const extra = {};
12     for (const k of Object.keys(entry)) {
13       if (!['id', 'seq', 'room', 'from', 'ts', 'kind', 'text', 'env', 'via'].includes(k)) extra[k] =
entry[k];
14     }
15     const res = q.insMsg.run(entry.id, room, seq, entry.from, entry.ts || Date.now(), entry.kind,
16       entry.kind === 'enc' ? null : (entry.text ?? null),
17       entry.kind === 'enc' ? JSON.stringify(entry.env) : null,
18       entry.via || null, Date.now(), Object.keys(extra).length ? JSON.stringify(extra) : null);
19     if (res.changes === 0) { // lost a race inside the same tx window
20       const row = q.getMsg.get(entry.id);
21       q.logDup.run(entry.id, room, entry.via || null, Date.now());
22       db.exec('COMMIT');
23       return { duplicate: true, rec: rowToEntry(row) };
24     }
25     db.exec('COMMIT');
26     feedCache.delete(room); // feed snapshot for this room is stale now
27     roomCount.set(room, countOf(room) + 1);
28     return { duplicate: false, rec: rowToEntry(q.getMsg.get(entry.id)) };
29   } catch (e) {
30     db.exec('ROLLBACK');
31     throw e;
32   }
33 };

```

8.4 Verification through the public URL

`scripts/dedup_test.py` verifies against the **database row count**, not against what a backend said:

Test	Scenario	Sent	Stored in DB	duplicate:true responses	Backends involved	Result
T1	same id, sequential retries	20	1	19	sys2, sys3, sys4	PASS
T2	same id, concurrent storm across backends	20	1	19	sys2, sys3, sys4	PASS
T3	send, drop connection, reconnect, re-send	2	1	1	sys3, sys4	PASS
T4	Idempotency-Key header	5	1	4	sys4	PASS
T5	control: distinct ids	20	20	0	sys3	PASS
P1	persistence across DB-service restart	49	49	1	via LB	PASS: count 48→49 across restart, probe readable=True, dup-of-old-id still rejected=True

Database after the test: 273769 messages, 10 duplicates rejected in total (sqlite (node:sqlite, WAL)).

The 20-way concurrent storm is the important one: twenty connections, spread by the balancer over all three backends, POST the same id at the same instant; exactly one row exists afterwards and all twenty responses carry the same seq. Test P1 shows that a duplicate of an id stored *before* a database restart is still rejected — the guard is durable, not in memory. The same property holds on the new route: the test suites post the same `/message` id repeatedly across two different backends and assert that exactly one is stored and that it appears exactly once in `/feed`.

9. Load Generator, Instrumentation and Test Suites

9.1 The load generator

`loadgen/loadgen.py` (Python 3 standard library, run from my own machine) drives the public URL. In its default `--api public` mode it uses exactly the two routes the evaluation generator will use — 60 % `POST /message`, 40 % `GET /feed`; `--api chat` drives the full authenticated application (10 % login, 55 % fetch, 30 % send, 5 % health) and is what the earlier measurements used.

The three kinds of variability the task asks for are all generator-side:

Requirement	Flag	Behaviour
variable number of users	<code>--concurrency N</code> , <code>--ramp "1:30,10:30,25:30,..."</code>	a fixed client count, or a schedule that raises and lowers it during the run
random/variable message length	<code>--msg-min 8</code> <code>--msg-max 240</code>	every message is a uniformly random length in characters, assembled from words so it looks like chat
random/variable interval	<code>--think-min 0</code> <code>--think-max 0.05</code> , or <code>--rate R</code>	a uniformly random pause between messages, or Poisson arrivals at a fixed offered rate

Modes: **closed loop** (N users in a request→response loop), **open loop** (arrivals at a fixed *offered* rate, independent of how slow responses get — this is what a throughput-vs-offered-load graph needs), and **ramp**. Per request it records latency, status, error class, `X-Backend-Id`, `X-LB-Active-Backends` and the duplicate flag; with `--poll-stats` it samples `/lb/stats` once a second, giving active backends, per-backend state, load index and CPU as a time series. The first seconds of every run are discarded as warm-up; percentiles are computed over the remaining sample. Every send carries a client-minted id and is retried with the same id on a transport error, exactly like the browser — so a failover run shows retries being *deduplicated*, not duplicated.

9.2 Measuring the utilisation of all four systems

The four systems are containers on one physical host, so `/proc/stat`, `uptime` and the load average are shared and describe the whole 120-core machine — all four report the same numbers and none of them is about that system. The per-system truth is each container's own cgroup v2 accounting, and `scripts/sysmetrics.py` samples it once a second over one persistent SSH connection per system:

File	Used for
<code>/sys/fs/cgroup/cpu.max</code>	quota / period → how many CPUs this system may use (here: exactly 1)
<code>/sys/fs/cgroup/cpu.stat</code>	<code>usage_usec</code> → CPU consumed; the delta over the interval is the utilisation
<code>/sys/fs/cgroup/memory.current</code> , <code>memory.max</code>	memory in use against the 512 MB allowance

`cpu_pct` is therefore a percentage of *that system's own CPU allowance*: 100 % means the container is using the whole CPU it was given. It is the same quantity the balancer reads out of `/health`, so the plots and the routing decisions are measured in the same units. Every experiment starts one of these samplers alongside the load generator, which is where the utilisation figures in §11.3 come from.

9.3 Automated test suites

Suite	Assertions	Covers
node <code>app/tests/smoke.js</code>	60	registration, login, sessions, encrypted rooms, message ids, dedup races across two backends, persistence across a database restart, self-registration and graceful deregister, and the /message + /feed routes : both body encodings, a missing <code>msg</code> , an id repeated across two different backends, <code>/feed</code> totals, truncation and cross-backend read-after-write

Suite	Assertions	Covers
python3 lb/test_lb.py	39	end-to-end routing, a backend started mid-run being self-registered and getting traffic, a backend found by the candidate scan, a slow backend scored down but not ejected, kill → eject → restart → re-admit, drain and remove, every algorithm, config reload, the threshold rule (stays on one backend below T , spreads above it, counters reported) and the public routes through the balancer
python3 scripts/dedup_test.py	5 scenarios	duplicate prevention against the live cluster through the public URL

All of them pass; §15 carries the captures.

9.4 The experiment matrix

`scripts/run_experiments_v2.sh` runs everything below against the public URL with a utilisation sampler attached, writes one JSON per run into `results/raw/` and appends a line to the run manifest.

Experiment	What varies	Held constant
THR threshold sweep	<code>switch_threshold</code> 0.15 ... 1.00, 2 reps, at 60 clients and at 10 clients	3 backends, 45 s per run
PALGO algorithm comparison	threshold vs round_robin vs least_connections vs adaptive, 2 reps	3 backends, 60 clients
P1/P2/P3 response time vs load	clients 1, 5, 10, 25, 50, 100, 200 × backends 1, 2, 3, 2 reps	the chosen threshold; the three backend counts are interleaved in shuffled order inside every level so the comparison shares the same host conditions
PO throughput vs offered load	offered 50, 100, 200, 300, 400 req/s (Poisson) × backends 1, 3	as above
PUTIL utilisation ramp	one 270 s run: 1 → 10 → 25 → 50 → 100 → 200 → 25 clients	3 backends
PFAIL failure / recovery	150 s at 60 clients; sys3 SIGKILLED at 45 s, restarted at 90 s	3 backends

The number of backends is always changed the *dynamic* way — by starting or stopping the backend process on the system (`scripts/scale.sh sysN up|down|kill`) — never by editing the balancer's configuration; the balancer must notice by itself.

10. Choosing the Optimal Threshold

"Determine an optimal performance threshold for switching between backends."

The threshold is chosen by measurement, not by taste. `scripts/run_experiments_v2.sh` `THR` sweeps `switch_threshold` from 0.15 to 1.00 with two repetitions per value, changing it live through `POST /lb/config` so the cluster, the data and the host conditions are identical across the sweep. `scripts/pick_threshold.py` then scores each value on what the evaluation generator will actually see, aggregating repetitions by **median** so one noisy run on a shared host cannot decide the answer:

$$\text{cost}(T) = \text{p95}(T) / \text{best p95} + 0.6 \times (\text{best throughput} - \text{throughput}(T)) / \text{best throughput} + 100 \times \text{error rate}(T)$$

p95 is the headline number; throughput is weighted a little lower because every value of `T` is offered the same load; any error rate at all is disqualifying. The winner is written to `results/processed/optimal_threshold.txt` and used by every later experiment and by the deployed configuration.

10.1 The sweep at 60 concurrent clients

Threshold sweep — 60 concurrent users on the public routes, 3 backends, median of the repetitions.

switch threshold T	reps	p50 (ms)	p95 (ms)	p99 (ms)	throughput (req/s)	errors (%)	busiest backend share	sys1 CPU (%)	backend CPU (%)
0.15	2	208	471	621	235.3	0.00	36 %	78	50
0.30	2	205	476	608	234.6	0.00	36 %	78	48
0.45	2	214	500	731	223.1	0.00	37 %	77	46
0.55	2	198	489	634	236.3	0.00	39 %	75	48
0.70 ← chosen	2	186	461	595	257.4	0.00	36 %	79	49
0.85	2	203	484	659	239.1	0.00	35 %	79	48
1.00	2	209	497	646	229.8	0.00	38 %	76	46




Figure 8 — Threshold sweep. Left and centre: 60 clients, whiskers show the min/max over the repetitions. Right: the same sweep at 10 clients, where the threshold decides whether traffic concentrates or spreads. Dotted line = the chosen `T`.

T = 0.70 wins on both metrics — the lowest median p95 (461 ms) and the highest throughput (257 req/s), roughly 7 % better than the worst value in the sweep on either. It is worth being honest about the size of that margin: the whole sweep spans 461–500 ms of p95 and 223–257 req/s, and the min/max whiskers on the figure show that the repetitions of a single value overlap the neighbouring values. At this load the choice of threshold is a second-order effect, and the reason is visible in the rightmost column of the table: at 60 clients **every** backend crosses **every** threshold within a fraction of a second, so the busiest-backend share sits at 35–39 % whatever `T` is, and what the sweep is really comparing is the quality of the tie-break rather than the threshold itself.

10.2 The same sweep at 10 concurrent clients — where the threshold actually bites

(pending)

At a load the cluster can comfortably absorb, the threshold decides something real: whether traffic concentrates on one warm backend or is spread across three. A high `T` keeps the cluster pinned; a low `T` spreads it immediately. Concentration is not automatically worse — a pinned backend keeps its connection pool and page cache hot — but it leaves no headroom for a burst, and it is why a value in the middle of the range is the right answer rather than either extreme.

10.3 The chosen configuration

```
{ "algorithm": "threshold", "switch_threshold": 0.70,  
  "inflight_cap": 24, "rt_cap_ms": 250, "explore_pct": 5 }
```

`T = 0.70` means: *keep using the current backend until it is running at 70 % of a CPU, or holding about 17 concurrent requests, or answering in more than 175 ms on average — then move.* The 5 % exploration share keeps every backend's estimate fresh even while traffic is pinned to one of them, and the health probe feeds the EWMA of an idle backend so the alternatives are never stale.

11. Results — Response Time, Throughput and System Utilisation

All of the following was produced by my own load generator against `http://10.1.75.53:3269`, using only `/message` and `/feed`, with random message lengths (8–240 characters) and random pauses (0–50 ms) between messages, and with the utilisation of all four systems sampled every second.

11.1 Response time vs load

Response time, throughput and per-system CPU against offered load on `/message` and `/feed`.

backends	users	throughput (req/s)	p50 (ms)	p95 (ms)	p99 (ms)	successful	failed	errors (%)	sys1 CPU	sys2 CPU	sys3 CPU	sys4 CPU
1	1	15.9	16	99	162	588	0	0.00	4 %	8 %	3 %	1 %
1	5	81.4	21	112	219	3013	0	0.00	24 %	39 %	13 %	1 %
1	10	107.7	45	147	323	3983	0	0.00	35 %	56 %	16 %	1 %
1	25	187.0	87	266	388	6919	0	0.00	55 %	75 %	16 %	1 %
1	50	181.0	235	444	616	6698	0	0.00	64 %	87 %	18 %	1 %
1	100	201.4	464	772	1302	7451	0	0.00	67 %	91 %	18 %	1 %
1	200	189.1	740	1281	9327	6997	0	0.00	58 %	94 %	18 %	1 %
2	1	14.4	24	142	248	1067	0	0.00	6 %	6 %	9 %	1 %
2	5	79.6	22	108	232	2944	0	0.00	22 %	27 %	32 %	1 %
2	10	140.0	26	125	304	5181	0	0.00	37 %	40 %	42 %	1 %
2	25	204.5	77	251	387	7566	0	0.00	63 %	51 %	59 %	1 %
2	50	66.4	152	443	652	2457	8	0.33	69 %	56 %	63 %	1 %
2	100	233.1	375	783	1041	8623	0	0.00	79 %	65 %	80 %	1 %
2	200	306.4	579	1186	1455	11337	0	0.00	82 %	67 %	84 %	1 %
3	1	16.7	15	96	153	619	0	0.00	4 %	9 %	6 %	3 %
3	5	82.0	21	106	198	3033	0	0.00	17 %	8 %	35 %	21 %
3	10	126.3	34	134	301	4674	0	0.00	36 %	21 %	43 %	33 %
3	25	203.1	79	246	373	7513	0	0.00	68 %	42 %	52 %	39 %
3	50	258.6	140	399	506	9570	0	0.00	73 %	44 %	58 %	48 %
3	100	225.4	388	820	1059	8340	0	0.00	72 %	43 %	61 %	48 %
3	200	229.5	805	1545	1898	8491	0	0.00	71 %	49 %	64 %	48 %




Figure 9 — Response time vs load on `/message` and `/feed` (p50 left, p95 right), 1 / 2 / 3 backends, log-log.




Figure 10 — Throughput vs load, and the effect of adding backends.

The shape is the textbook one. Up to about 25 clients the system is **latency-bound**: response time is flat and throughput rises almost linearly with the offered load. Past that it becomes **capacity-bound**: throughput flattens and every additional client simply adds queueing time, so the median rises roughly in proportion to the client count. With three backends the knee is at 50 clients and 259 req/s; a single backend reaches its knee earlier and lower.

Two features of the curves are worth reading carefully. First, **the p95 at one client is six times the p50** — that is not noise but the request mix: a `/message` costs about 15 ms while a `/feed` carries 54 KB across the campus link, so the median is a write and the tail is a read. Second, adding backends helps least at very low load, where nothing is queueing anyway, and most in the middle of the range, where the extra CPU converts directly into throughput.

11.2 Throughput vs offered load

Open-loop throughput: Poisson arrivals at a fixed offered load, independent of how slow the answers get.

backends	offered load (req/s)	achieved throughput (req/s)	p95 (ms)	errors (%)	arrivals dropped
1	50	46.8	214	0.00	0
1	100	84.6	456	0.00	0
1	200	130.5	8667	1.35	0
1	300	114.5	20724	44.50	2941
1	400	106.9	19495	61.53	5971
3	50	47.0	222	0.00	0
3	100	84.0	323	0.00	0
3	200	153.1	3096	0.00	0
3	300	115.9	20450	42.05	2606
3	400	111.0	18842	60.25	5806




Figure 11 — Throughput vs offered load (open loop, Poisson arrivals); dashed = ideal.

The open-loop sweep offers a fixed arrival rate regardless of how slow the answers get, which is what exposes overload. Both configurations track the ideal line to 100 req/s. At 200 req/s offered, the single backend has already collapsed — 8.7 s at p95 and the first failures — while three backends still answer in 3.1 s with **zero** failed requests, which is the clearest single comparison in the report. Beyond 300 req/s offered both configurations fail hard; that is genuine overload of a three-core cluster, and the "arrivals dropped" column records where my generator itself refused to add more in-flight requests rather than measure its own thread pool.

11.3 System utilisation of all four systems

Mean CPU of each of the four systems while the offered load rises (3 backends). Each system is a container with a one-CPU quota, so 100 % is its whole allowance.

backends	users	throughput (req/s)	sys1 — load balancer	sys2 — backend	sys3 — backend + DB	sys4 — backend
3	1	16.7	4 %	9 %	6 %	3 %
3	5	82.0	17 %	8 %	35 %	21 %
3	10	126.3	36 %	21 %	43 %	33 %
3	25	203.1	68 %	42 %	52 %	39 %
3	50	258.6	73 %	44 %	58 %	48 %
3	100	225.4	72 %	43 %	61 %	48 %
3	200	229.5	71 %	49 %	64 %	48 %




Figure 12 — Utilisation of all four systems during the 1 → 200 client ramp, with the client count and p95 response time underneath. 100 % is one whole CPU, which is each container's quota.




Figure 13 — Mean CPU of each of the four systems against offered load, 3 backends.

This is the plot the updated task asks for, and it is where the system's real limit shows up.

- **sys1, the load balancer, is the busiest system in the cluster.** It climbs to about 70 % of its CPU allowance by 25 clients and peaks near 95 % during the ramp, while each backend sits between 40 % and 65 %. Every client connection is terminated here and every byte of every `/feed` response is copied through this one Python process on a one-CPU container, so it saturates first.
- **sys3 runs hotter than sys2 and sys4** at the same offered load, because it also carries the shared database. The balancer sees that in the load index and sends it less traffic — 28 % of requests against

35 % and 37 % for the other two in the algorithm comparison. This is exactly the behaviour the threshold rule exists for, demonstrated by accident of the deployment rather than by an artificial CPU hog.

- **The backends are not the bottleneck** past 25 clients: their utilisation is flat from 50 clients upward while response time keeps climbing, which is the signature of the queue being somewhere else.
- Memory was never a constraint: each backend held 230–260 MB of its 512 MB, the balancer under 30 MB, and the database service about 100 MB plus SQLite's page cache.

11.4 The threshold rule against the classic algorithms

The chosen threshold rule against the classic algorithms, same load, same cluster.

algorithm	reps	throughput (req/s)	p50 (ms)	p95 (ms)	p99 (ms)	errors (%)	busiest backend	backend CPU spread	traffic split
threshold	2	260.0	184	447	565	0.00	38 %	17 pp	sys2 35% / sys3 28% / sys4 37%
least_connections	2	240.8	197	471	620	0.00	34 %	19 pp	sys2 32% / sys3 34% / sys4 34%
adaptive	2	233.9	202	489	679	0.00	34 %	18 pp	sys2 34% / sys3 32% / sys4 34%
round_robin	2	220.6	224	494	646	0.00	33 %	21 pp	sys2 33% / sys3 33% / sys4 33%




Figure 14 — The threshold rule against round robin, least connections and the adaptive score, same load and same cluster.

Same load, same cluster, same three backends, only the selection rule changes:

- **threshold (T = 0.70) is the fastest and the highest-throughput rule** — 260 req/s at 447 ms p95, which is **+18 % throughput and –10 % p95 against fixed round robin** (221 req/s, 494 ms).
- Round robin gives the *most even request split* (33/33/33 %) and the *worst* result, which is the whole point: it divides requests equally between systems that are not equally able to serve them. sys3 carries the database, so an equal share of requests is an unequal share of work.
- The load-aware rules all beat it, and they all send sys3 less traffic. The threshold rule is the most decisive about it (28 % to sys3) and wins by the largest margin; `least_connections` and the `adaptive` score land in between.
- No configuration produced a single failed request at this load.

12. Results — Dynamic Scaling, Failure Recovery and Persistence

These are the requirements carried over from the original task. §12.1–12.3 were measured on the authenticated chat API with the adaptive algorithm, before the database moved off sys1; §12.4 repeats the

failure test on the required public routes with the threshold algorithm, and §12.5 covers persistence and duplicate prevention on the live cluster.

12.1 Dynamic scaling: performance after each backend is added

The system started with **one** backend (sys2). Load was raised in steps (25 → 50 → 100 users) until the single backend was saturated; then `bash scripts/scale.sh sys3 up` and later `... sys4 up` were executed on the other systems *while the load generator kept running*. The balancer admitted each new backend by itself (registration heartbeat and candidate scan both fired) and the load generator's per-second samples of `/lb/stats` show exactly when.

Phase	Window (s)	Users	Active backends	Throughput (req/s)	p50 ms	p95 ms	Errors	Share sys2 / sys3 / sys4
A: 1 backend, 25 users	5–40	25	1	95.5	75	2089	0	100% / 0% / 0%
B: 1 backend, 50 users	45–80	50	1	105.2	75	4482	0	100% / 0% / 0%
C: 1 backend, 100 users (saturated)	85–128	100	1	98.6	61	9028	0	100% / 0% / 0%
D: sys3 added → 2 backends	136–184	100	2	189.5	100	4172	0	52% / 47% / 1%
E: sys4 added → 3 backends	192–260	100	3	285.2	102	2432	0	34% / 34% / 32%
F: 3 backends, 200 users	265–300	200	3	291.1	133	5710	0	34% / 33% / 33%




Figure 4 — Dynamic scaling timeline: throughput, response time, active backends and per-backend share; dashed lines = backend added.




Figure 5 — Effect of adding each backend, per phase of the scaling run.

How performance changed after each backend was added. With one backend the system was already at its ceiling at 25 users (~95 req/s, p95 ≈ 2.1 s): raising the load to 50 and then 100 users brought *no* throughput (105 → 99 req/s) and only pushed queueing delay up — p95 climbed to 4.5 s and then 9 s, the classic closed-loop saturation signature (the scaling figure, phases A–C). At ~128 s `scale.sh sys3 up` was run on sys3; its first registration heartbeat and the balancer's candidate scan both fired within 5 s, the `active backends` line steps to 2 and — because a never-sampled backend scores 0 — sys3 was carrying real traffic in the very next second (share panel: 52 / 47 %). Throughput **doubled to ~190 req/s** and p95 halved to 4.2 s with the *same* 100 users. Adding sys4 at ~184 s repeated the effect: **~285 req/s**, p95 2.4 s, a near-perfect 34 / 34 / 32 % split. Phase F then doubled the load to 200 users: throughput held at ~291 req/s (the three cores are now the ceiling) and p95 rose again to 5.7 s. Scaling from one to three backends therefore gave **2.9× throughput** (95–105 → 285–291 req/s) — near-linear, as expected for CPU-bound backends behind a balancer whose own cost is small — and no request failed during either addition (0 errors in 300 s). The `added` events in the dashboard (screenshot §15) are the balancer's own record of the two admissions.

12.2 Backend failure and recovery under load

Phase	Window (s)	Users	Active backends	Throughput (req/s)	p50 ms	p95 ms	Errors	Share sys2 / sys3 / sys4
healthy, 3 backends	5–29	100	3	280.4	125	2357	0	31% / 33% / 35%
sys3 killed → detection window	29–39	100	3	250.0	88	2737	36	34% / 26% / 39%
2 backends carry the load	39–90	100	2	184.1	94	4421	0	47% / 0% / 53%
sys3 restarted → re-admitted	90–100	100	3	225.0	182	1935	0	32% / 30% / 39%
3 backends again	100–150	100	3	279.8	96	2610	0	33% / 33% / 34%

Whole run: 34927 requests, 36 failed (0.103 %), 2 sends retried with the same id, 0 of those answered `duplicate:true` (stored once).




Figure 6 — Failure and recovery timeline under 100 users (sys3 SIGKILLED at ~37 s, restarted at ~92 s; dashed = LB ejected / re-admitted).

Failure. sys3 was SIGKILLED (no graceful deregister — the worst case) at ~37 s while 100 users were active. Requests that were in flight on sys3 at that instant, plus the few that the balancer sent before its passive check fired, failed: **36 requests out of 34 927 (0.10 %)**, all inside a two-second window. The first connection refusal ejected sys3 immediately (ejected ... passive: connect/proxy failure in the event log — no waiting for a health-check interval), and the active probe confirmed it. Throughput fell from ~280 to ~184 req/s and p95 rose from 2.4 s to 4.4 s — exactly the two-thirds capacity you expect with two of three cores left — and the share panel shows the load redistributed 47 / 53 % over sys2 and sys4 without any further errors. The load generator's sends that hit the outage were retried with the same message id; they were stored once (the first attempt never reached the database, so the retry is a fresh insert — dedup is verified separately in §8.4).

Recovery. scale.sh sys3 up at ~90 s: the backend registered itself, the balancer re-admitted it after two successful probes (readmitted event, active backends back to 3) and, scoring 0 as a fresh backend, it took its third of the traffic at once; throughput returned to ~280 req/s within ten seconds. Detection-to-ejection time is bounded by max(passive: first failed connect, active: $2 \times 3 \text{ s} + \text{timeout}$) — in this run it was effectively instantaneous.

12.3 Dynamic selection vs fixed rotation with one loaded backend

To show that selection really follows load, a CPU-burning process (`python3 -c 'while True: pass'`) was started on sys3 — on a 1-core container it takes roughly half of the core away from the backend. The same 50-user load was then run under each algorithm.

Algorithm	sys3 CPU hog	Throughput (req/s)	p50 ms	p95 ms	p99 ms	Errors	Share sys2 / sys3 / sys4
adaptive	yes	247.0	89.06	1135.83	1841.14	0	37.4% / 25.0% / 37.6%
round_robin	yes	201.81	38.8	492.81	5847.76	0	33.4% / 33.1% / 33.4%
least_connections	yes	257.33	80.31	1252.91	1898.96	0	35.5% / 27.5% / 36.9%
adaptive	no	271.83	86.13	1021.12	1404.22	0	32.1% / 35.6% / 32.3%
round_robin	no	275.46	65.35	1109.0	2149.48	0	33.3% / 33.3% / 33.3%




Figure 7 — Adaptive vs round robin vs least connections when sys3 is CPU-loaded.

With sys3 losing half its core, the balancer saw it at once (`sys_cpu_pct` 100 %, DEGRADED, EWMA rising) and the **adaptive algorithm cut sys3's share from a third to 25 %**, which is worth **22 % more throughput than round robin (247 vs 202 req/s)** at the same 50 users. Round robin kept feeding sys3 one third and paid with a **p99 of 5.8 s** (requests queued on the loaded backend). Without the hog the two rules are indistinguishable (272 vs 275 req/s): the dynamic rule costs nothing when there is nothing to react to and pays when one backend is quietly slower.

12.4 Failure and recovery on the required routes

The same experiment repeated on `/message` and `/feed` with the threshold algorithm: 60 clients for 150 s, sys3 SIGKILLED at 45 s and restarted at 90 s.




Figure 15 — Failure and recovery on the public routes: per-backend throughput, response time, active backends and failed requests. sys3 was SIGKILLED at 45 s and restarted at 90 s.

The balancer ejects sys3 on the first refused connection, the other two absorb its share, and the `active backends` line steps $3 \rightarrow 2 \rightarrow 3$. Because sys3 also hosts the database, this run kills a *backend* on the database's system while the database itself keeps serving — the two are separate processes with separate lifetimes, which is why `scale.sh sys3 kill` does not take the cluster down.

12.5 Persistence and duplicate prevention on the allotted systems

Persistence. The database file held 273 769 messages, 719 users and 7 rooms when the original experiments finished (`/stats` , terminal capture §15) — every load-generator message of every run, the migrated Assignment-5 history, and the demo conversations, all served identically by whichever backend the balancer picks. `dedup_test.py --restart-db` restarted the database service in the middle of the test (test P1): the row count was unchanged across the restart (48 $\rightarrow$ 49 with the probe message), the probe message was readable through a different backend afterwards, and re-sending an id stored *before* the restart was still answered `duplicate:true` — the guard lives in the database file, not in any process's memory. The backends reconnected to the database's firehose by themselves (`db_firehose: true` in `/health`). The smoke test exercises the same restart locally on every run.

Duplicate prevention. All five duplicate scenarios in §8.4 passed through the public URL against three live backends, and the `dedup_log` table holds an audit row for every rejection the database itself had to make (10 in total; most duplicates never reach it because the backends' LRU answers them first). A SQL check on the live file — `SELECT id FROM messages GROUP BY id HAVING COUNT(*) > 1` — returns no rows (terminal capture §15).

13. Bottleneck Analysis — Why the Database Moved to sys3

This is the measurement that changed the deployment, and it is the most useful thing the utilisation sampler produced.

The symptom. With the database service next to the balancer on sys1, 60 concurrent clients on the public routes produced a hard ceiling: **165 req/s, 324 ms median, 693 ms p95**, and it would not move whatever the threshold was set to. The backends were only 36 % busy, so the queue was not there.

The measurement. `top` on sys1 showed the balancer at ~60 % of a CPU and the database service at ~38 % — together ~98 % of the container's single-CPU quota. The decisive number came from `/sys/fs/cgroup/cpu.stat`, sampled over an 8-second window during a run:

	value
scheduling periods elapsed	80
periods in which the container was throttled	41
CPU time consumed	6.55 s (82 % of one core)
time tasks spent stopped by the quota	9.06 s

sys1 was being CFS-throttled in **half of all 100 ms periods**. When a container exhausts its quota inside a period, every task in it is stopped until the next period begins — which is exactly the 300 ms of unexplained queueing the response times showed. The two most latency-critical processes in the system were competing for one core and taking turns being frozen.

The check that it was not something else. Three alternatives were ruled out by measurement rather than argument. Running the load generator *inside* the lab (from sys4, over the bridge network) produced the same 163 req/s, so the campus link was not the limit. The balancer was benchmarked in isolation against a trivial backend and sustained **11 000–12 000 req/s with 54 KB bodies**, and its throughput was flat from 12 to 120 concurrent connections, so neither the proxy code nor the thread-per-connection model was at fault. A CPU benchmark put sys1's core at 3.5× slower than the machine the balancer was benchmarked on — nowhere near enough to explain a 70× gap.

The fix and its effect. The database service, its data file and its WAL were moved to sys3 (`bash scripts/deploy.sh movedb sys1 sys3`, WAL checkpointed first so no row was lost), leaving sys1 to the balancer alone. At identical load:

	before (DB on sys1)	after (DB on sys3)	change
throughput	165 req/s	233 req/s	+41 %
median response time	324 ms	230 ms	−29 %
p95 response time	693 ms	529 ms	−24 %
sys1 throttled periods (per 80)	41	30	−27 %

Three smaller changes were made in the same pass, each for a measured reason: the balancer's access log is now accumulated and flushed once a second instead of costing a `write` syscall per request; the pooled upstream connections keep their buffered reader instead of allocating a new one per request; and the sockets on both sides get 256 KB buffers, because a 54 KB body arriving over a container bridge otherwise turns one logical read into a dozen `recv` calls. On the database side the room's row count is maintained incrementally instead of running `COUNT(*)` on every `/feed`, and both the database and the backends cap how often the same feed window is re-serialised. Together these took the same load from 165 to 238 req/s.

What it costs and why that is acceptable. sys3 now carries a backend *and* the database, so it is the slowest of the three backends. That is precisely the asymmetry a performance-based balancer is supposed to handle,

and §11.4 shows it doing so — sys3 receives 28 % of requests where round robin would give it 33 %. The alternative, an even split across unequal systems, is measurably worse.

What is still the limit. sys1 remains the busiest system at ~70–95 % of its single CPU, so the cluster ceiling is now the balancer's own CPU rather than a scheduling artefact. Adding a fourth backend would not help; giving sys1 a second core, or running two balancer processes behind `SO_REUSEPORT` on a multi-core system, would.

14. Analysis and Discussion

The threshold rule does what the task asks and is measurably better than round robin. At the same 60-client load on three backends it delivers 260 req/s at 447 ms p95 against round robin's 221 req/s at 494 ms — **+18 % throughput, -10 % p95**. The reason is visible in the traffic split: round robin gives all three systems an equal share of *requests*, but sys3 also runs the database, so an equal share of requests is an unequal share of work. The threshold rule reads the load index, sees sys3 sitting higher, and gives it 28 % instead of 33 %. Where all three backends are genuinely equal the two rules are indistinguishable, which is the correct behaviour: a dynamic rule should cost nothing when there is nothing to react to.

The threshold itself matters less than the deployment. The honest reading of the sweep in §10 is that any value between 0.15 and 1.00 lands within about 7 % of the best, and the repetitions overlap. Moving the database off the balancer's system was worth 41 % — six times the entire spread of the threshold sweep. Tuning is real but second-order; finding the actual queue is first-order. The sweep was still worth running, because it is what proved that.

The bottleneck is the balancer, not the backends. Every client connection terminates on sys1 and every byte of every response is copied through one Python process on a one-CPU container. Past 25 clients the backends flatten at 40–65 % utilisation while sys1 climbs to 70–95 % and response times keep rising — the queue is in front of the backends, not in them. This is why the "effect of adding a backend" is strong from one to two backends and weaker from two to three, and why a fourth backend would not help.

Scaling is still real, and it is what protects the system under overload. The clearest comparison in the report is the open-loop sweep at 200 req/s offered: one backend is at 8.7 s p95 and failing, three backends answer in 3.1 s with zero failures. Under closed-loop load three backends reach their knee at 50 clients and 259 req/s where one backend is already saturated at 25.

Persistence and idempotency are properties of the data model, not of luck. The client-minted id is the primary key; retries, reconnections, and twenty concurrent copies of the same message through three different backends produce exactly one row, and a duplicate of an id stored *before* a database restart is still rejected afterwards. Answering duplicates with `200 duplicate:true` rather than an error is deliberate: a client that receives an error would retry again. The same guarantee covers the new `/message` route, and the test suites assert it across two different backends.

Measurement discipline. The lab host is shared and multi-tenant, and single runs vary: one repetition of the same configuration came back at 30 req/s where its twin gave 81. Everything reported here is therefore a median over repetitions, the backend-count comparisons are interleaved in shuffled order inside every load level so they share the same minutes, and the threshold figure shows the min/max of the repetitions as whiskers rather than hiding them.

Limits and future work. The database service is a single point of failure and, eventually, a bottleneck; the standard next steps are a replicated store (PostgreSQL with streaming replication, or leader/follower SQLite via Litestream) and backend-local write queues. The balancer's DEGRADED rule still uses fixed thresholds (2 s probe, 95 % CPU) where an adaptive bound derived from the EWMA's own history would be more robust. The open-loop generator should use an asynchronous client to offer more than ~300 arrivals/s from one process. TLS termination at the balancer and a second balancer behind a shared virtual IP would remove the last single point of failure.

15. Screenshots and Evidence

Live LB dashboard — three backends UP, adaptive scores, event log

Dynamic load balancer — sys1:3269

algorithm **adaptive** · 3 active backend(s) of 3 · uptime 619s · 97117 requests · 38 errors · auto-refresh 2s

backend	addr	source	state	score ↓	EWMA ms	probe ms	cpu %	load1	in-flight	requests	errors	p50/p95 ms	share
sys2	172.17.0.71:3000	config	DEGRADED	35595.2	278.1	104.4	100	3.77	31	44096	0	81 / 2947	45.8%
sys3	172.17.0.72:3000	scan	DEGRADED	53034.0	736.6	25.6	100	3.77	17	25928	38	94 / 2575	26.9%
sys4	172.17.0.73:3001	scan	DEGRADED	36923.5	426.2	66.9	96.9	3.77	21	26184	0	74 / 2758	27.2%

Recent events

- 07:18:25 degraded sys3 probe 1762 ms, cpu 100% (active=3)
- 07:18:20 readmitted sys3 probe 6 ms (active=3)
- 07:17:24 ejected sys3 passive: connect/proxy failure (active=2)
- 07:16:50 degraded sys4 probe 80 ms, cpu 96% (active=3)
- 07:16:50 degraded sys3 probe 151 ms, cpu 100% (active=3)
- 07:16:50 degraded sys2 probe 251 ms, cpu 96% (active=3)
- 07:16:43 config - {"algorithm": "adaptive"} (active=3)
- 07:16:38 recovered sys4 probe 5 ms (active=3)
- 07:16:38 recovered sys3 probe 5 ms (active=3)
- 07:16:38 recovered sys2 probe 5 ms (active=3)
- 07:14:39 degraded sys4 probe 1663 ms, cpu 100% (active=3)
- 07:14:34 readmitted sys4 probe 10 ms (active=3)

JSON: [/lb/stats](#) · [/lb/events](#) · reload config: POST /lb/reload · register: POST /lb/register

Dashboard during the scaling run — sys3 and sys4 admitted while load was running

Dynamic load balancer — sys1:3269

algorithm **adaptive** · 3 active backend(s) of 3 · uptime 422s · 50496 requests · 0 errors · auto-refresh 2s


backend	addr	source	state	score ↓	EWMA ms	probe ms	cpu %	load1	in-flight	requests	errors	p50/p95 ms	share
sys2	172.17.0.71:3000	config	DEGRADED	51223.0	312.3	6.6	100	2.77	40	27627	0	88 / 2353	55.3%
sys3	172.17.0.72:3000	scan	DEGRADED	25109.3	190.4	137.6	99.8	2.87	32	13869	0	75 / 2462	27.8%
sys4	172.17.0.73:3001	scan	DEGRADED	111245.2	1071.3	61.4	99.7	2.87	25	8475	0	81 / 2584	17.0%

Recent events


- 07:14:39 degraded sys4 probe 1663 ms, cpu 100% (active=3)
- 07:14:34 readmitted sys4 probe 10 ms (active=3)
- 07:13:43 degraded sys3 probe 437 ms, cpu 98% (active=2)
- 07:13:39 readmitted sys3 probe 6 ms (active=2)
- 07:11:17 ejected sys4 drained backend went away (active=1)
- 07:11:17 ejected sys3 drained backend went away (active=1)
- 07:11:17 degraded sys2 probe 14 ms, cpu 100% (active=1)
- 07:11:14 draining sys4 deregister request (active=1)
- 07:11:13 draining sys3 deregister request (active=2)
- 07:11:12 config - {"algorithm": "adaptive"} (active=3)
- 07:09:23 recovered sys2 probe 9 ms (active=3)
- 07:09:19 degraded sys2 probe 7 ms, cpu 100% (active=3)

JSON: [/lb/stats](#) · [/lb/events](#) · reload config: POST /lb/reload · register: POST /lb/register

Chat login screen served through the load balancer



Conversation: message ids, seq numbers, serving backend, cluster panel



'Resend last' — same message id re-sent, server answers duplicate:true, nothing stored twice

rahul_demo

Sign out

new-room-name

+

assignment6-demo

dedup-test

general

loadtest

project-demo

public-room

secrets

assignment6-demo

3 stored

Lock room

rahul_demo · 12:50:53 · #1 · via sys2

Hello from Assignment 6 — this message has a client-minted UUID id and is stored once in SQLite

rahul_demo · 12:51:46 · #2 · via sys2

Second message — sent through the dynamic load balancer, delivered via WebSocket, persisted on sys1

rahul_demo · 12:53:19 · #3 · via sys2

Duplicate test: this message is re-sent with the SAME message id

already stored as #3 — duplicate suppressed (backend-lru)

Duplicate test: this message is re-sent with the SAME message id (resend with same id)

Cluster

3 active · adaptive

sys2DEGRADED

502ms

cpu

97.8%

sys3DEGRADED

386ms

cpu

99.4%

sys4DEGRADED

213ms

cpu

99.7%

Connected

sys2

Message... (Enter to send)

Send

Resend last

Dashboard with sys3 DOWN during the failure test

Dynamic load balancer — sys1:3269

algorithm adaptive · 2 active backend(s) of 3 · uptime 43s · 37 requests · 0 errors · auto-refresh 2s

backend	addr	source	state	score ↓	EWMA ms	probe ms	cpu %	load1	in-flight	requests	errors	p50/p95 ms	share
sys2	172.17.0.71:3000	config	UP	11.6	5.7	5.9	0.2	0.88	1	3	0	2 / 2	37.5%
sys3	172.17.0.72:3000	scan	DOWN	0.0	—	5.0	0.7	1.04	0	1	0	0 / 0	12.5%
sys4	172.17.0.73:3001	scan	UP	5.7	5.6	10.3	0.1	0.88	0	4	0	2 / 2	50.0%

Recent events

- 08:14:21 ejected sys3 passive: connect/proxy failure (active=2)
- 08:14:08 recovered sys3 probe 5 ms (active=3)
- 08:14:02 degraded sys3 probe 10 ms, cpu 100% (active=3)
- 08:13:50 added sys4 172.17.0.73:3001 via scan (active=3)
- 08:13:50 added sys3 172.17.0.72:3000 via scan (active=2)

JSON: [/lb/stats](#) · [/lb/events](#) · reload config: POST /lb/reload · register: POST /lb/register

Terminal evidence

01_sanity_check.txt

```
$ bash scripts/sanity_check.sh
== Assignment 6 sanity check Fri Sep  4 13:54:18 IST 2026 ==
[✓] LB v2 (public)      http://10.1.75.53:3269/lb/health {"status":"ok","service":"lb","version":"v2-
dynamic","active_backends":3}
[✓] app via LB          http://10.1.75.53:3269/health {"status":"ok","backend":"sys2","version":"v3-
assignment6","uptime":634,"connections":1,"db_firehose":true,"lo
algorithm=adaptive active_backends=3 total_requests=81757
- sys2 172.17.0.71:3000 UP      src=config ewma=37.2 score=74.9 cpu=0.8 reqs=28220
- sys3 172.17.0.72:3000 UP      src=scan ewma=6.1 score=6.1 cpu=0.4 reqs=24219
- sys4 172.17.0.73:3001 UP      src=scan ewma=6.1 score=6.2 cpu=0.8 reqs=28582
[✓] DB service sys1:5270 sqlite messages=273769 users=720 rooms=7 dups_rejected=10 subscribers=3
[✓] backend sys2 :3000 v3-assignment6 cpu 0.5 %
[✓] backend sys3 :3000 v3-assignment6 cpu 0.4 %
[✓] backend sys4 :3001 v3-assignment6 cpu 0.4 %
-- previous assignments --
[✓] Lab5 direct sys2 http://10.1.75.53:3270/health {"status":"ok","backend":"sys2","version":"v3-
assignment6","uptime":634,"connections":1,"db_firehose":true,"lo
[✓] Lab5 direct sys3 http://10.1.75.53:3271/health {"status":"ok","backend":"sys3","version":"v3-
assignment6","uptime":584,"connections":0,"db_firehose":true,"lo
[✓] Project Contour http://10.1.75.53:3272 (sys4:3000, untouched)
[✓] Lab5 state service sys1:5269 (still running, untouched)
```

03_db_stats_schema.txt

```
$ ssh lbsys1 'curl -s http://127.0.0.1:5270/stats'
{
  "service": "db",
  "engine": "sqlite (node:sqlite, WAL)",
  "file": "/home/student/assignment6/data/chat.sqlite",
  "db_bytes": 74354688,
  "users": 720,
  "rooms": 7,
  "messages": 273769,
  "duplicates_rejected_total": 10,
  "duplicates_rejected_since_boot": 0,
  "appended_since_boot": 0,
  "requests": 7,
  "subscribers": 3,
  "uptime": 49,
  "node": "v22.23.2"
}

$ ssh lbsys1 'python3 -c "import sqlite3; ...schema, counts, duplicate-id check..."'
CREATE TABLE dedup_log (
  n INTEGER PRIMARY KEY AUTOINCREMENT, id TEXT NOT NULL, room TEXT, via TEXT, at INTEGER NOT NULL);
CREATE TABLE messages (
  id TEXT PRIMARY KEY,
  room TEXT NOT NULL,
  seq INTEGER NOT NULL,
  sender TEXT NOT NULL,
  ts INTEGER NOT NULL,
  kind TEXT NOT NULL,
  text TEXT,
  env TEXT,
  via TEXT,
  stored_at INTEGER NOT NULL,
  extra TEXT,
  UNIQUE(room, seq)
);
CREATE TABLE rooms (id TEXT PRIMARY KEY, locked INTEGER NOT NULL DEFAULT 0, epoch INTEGER NOT NULL DEFAULT 0,
verifier TEXT, created INTEGER NOT NULL, created_by TEXT);
CREATE TABLE sessions (token TEXT PRIMARY KEY, username TEXT NOT NULL, created INTEGER NOT NULL, expires INTEGER NOT
NULL);
CREATE TABLE users (username TEXT PRIMARY KEY, hash TEXT NOT NULL, created INTEGER NOT NULL);
CREATE INDEX idx_messages_room_seq ON messages(room, seq);
CREATE INDEX idx_sessions_expires ON sessions(expires);

rows: {'messages': 273769, 'users': 720, 'rooms': 7, 'dedup_log': 10}
duplicate ids in messages (must be 0): 0
recent dedup_log: [('c904a223-fa2c-4120-9cfd-333ca229e693', 'dedup-test', 'sys4'), ('c904a223-fa2c-4120-9cfd-
333ca229e693', 'dedup-test', 'sys2'), ('c904a223-fa2c-4120-9cfd-333ca229e693', 'dedup-test', 'sys4')]
```

04_ssh_listeners.txt

```
$ bash scripts/ssh_all.sh 'hostname; hostname -I; ss -tlnp 2>/dev/null | grep -E ":(3000|3001|5270) "'
===== lbsys1 =====
[lbsys1] stu78_sys1
[lbsys1] 172.17.0.70
[lbsys1] LISTEN 0      511          0.0.0.0:5270    0.0.0.0:*      users:(("node",pid=255461,fd=21))
[lbsys1] LISTEN 0      512          0.0.0.0:3000    0.0.0.0:*      users:(("python3",pid=255529,fd=5))
[lbsys1] exit: 0
===== lbsys2 =====
[lbsys2] stu78_sys2
[lbsys2] 172.17.0.71
[lbsys2] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("node",pid=241117,fd=18))
[lbsys2] exit: 0
===== lbsys3 =====
[lbsys3] stu78_sys3
[lbsys3] 172.17.0.72
[lbsys3] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("node",pid=240499,fd=18))
[lbsys3] exit: 0
===== lbsys4 =====
[lbsys4] stu78_sys4
[lbsys4] 172.17.0.73
[lbsys4] LISTEN 0      511          0.0.0.0:3001    0.0.0.0:*      users:(("node",pid=259908,fd=18))
[lbsys4] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("nginx",pid=233485,fd=5),
("nginx",pid=231101,fd=5))
[lbsys4] exit: 0
```

16. Challenges Faced and How They Were Solved

1. **The cluster stalled at 165 req/s and nothing in the application explained it.** The backends were only 36 % busy. `cpu.stat` showed sys1 being CFS-throttled in half of all scheduling periods because the balancer and the database were sharing its single CPU. Moving the database to sys3 was worth 41 % throughput (§13) — and finding it needed per-container accounting, because `/proc` on these boxes describes the whole physical host.
2. **`/feed` returning literally every message made the benchmark measure the history, not the system.** At 1 412 messages the response was already 383 KB and 190 ms, growing linearly. Resolved by making the default a 200-message window that always reports the true total and whether it was truncated, with `?limit=all` for the complete history — bounded by default, nothing unreachable.
3. **A pure "least response time" rule herded all traffic onto one backend.** With sequential traffic every backend has zero in-flight requests, so the one with the lowest EWMA received everything. The same trap appears in the threshold rule when several requests cross the threshold at once and all pick the same replacement. Both are fixed with power-of-two-choices (§6.1).
4. **The load signal was wrong for a loaded box.** A CPU-burning neighbour *reduces* the Node process's own CPU share, which made a loaded backend look idle. The backend now also reports the container's cgroup CPU against its quota, which sees every tenant, and the balancer takes the maximum of the two.
5. **sys3 had Node 20, which has no SQLite module, and no internet access to install one.** The user-local Node 22 tree was copied from sys1 over SSH; the system Node is untouched and the previous assignment's services keep using it.
6. **Port 3000 on sys4 was already taken** by my course project. Only the balancer needs a public port, so the sys4 backend listens on 3001 on the private network and nothing of the project was touched.
7. **The evaluation generator's request format is unspecified.** `/message` accepts JSON, form-encoded bodies, raw text and query parameters, `GET` as well as `POST`, and eight spellings of each field name, so a reasonable client cannot fail to be understood.

17. Conclusion

The secure group chat now runs on the three allotted systems behind a load balancer that clients reach at a single URL, `http://10.1.75.53:3269`, through the two required routes `/message` and `/feed`. The balancer selects backends by **measured load** — a threshold on an index combining container CPU, queue depth and

EWMA response time — switching away from the current backend the moment it crosses $T = 0.70$, a value chosen by a two-level sweep rather than by assumption. It monitors health in four states, ejects an unavailable backend on the first refused connection while keeping a merely slow one in service, and admits a backend started while the application is running within one heartbeat.

Measured on the allotted systems with my own load generator, driving only the required routes with random message lengths and random intervals: the threshold rule delivers **260 req/s at 447 ms p95** against fixed round robin's **221 req/s at 494 ms**; three backends reach their knee at 50 clients and 259 req/s where one saturates at 25; at 200 req/s of offered open-loop load a single backend is at 8.7 s p95 and failing while three backends answer in 3.1 s with **zero** failures; a SIGKILLED backend is ejected within one request and back in rotation seconds after restarting; and the utilisation trace of all four systems shows the balancer's own CPU, not the backends, as the remaining ceiling. Diagnosing that ceiling — CFS throttling on a shared one-CPU container — and moving the database off the balancer's system was worth **+41 %** throughput, more than the entire span of the threshold sweep.

All backends share one persistent SQLite database in which every message has a unique id and the database itself guarantees that a retried, reconnected or concurrently duplicated send is stored once; every duplicate test through the public URL leaves exactly one row. **Nothing was removed to reach these numbers**: the two required routes are additions, and registration, script login, sessions, end-to-end-encrypted rooms and the full `/api/*` surface all still work through the same URL and are still covered by the 60-assertion backend suite and the 39-assertion balancer suite, both of which pass. The previous assignment's URLs and direct ports keep working, its files on the lab systems are untouched, and the whole system can be redeployed, re-measured and demonstrated from the scripts in the repository.

18. Appendix — The Same Sweeps on the Authenticated Chat API

The measurements below drive `/api/*` — register, script login, fetch history, send — rather than the two public routes, and were taken with the adaptive algorithm before the database moved off sys1. They are kept because they are the evidence for the original task and because they show the same system under a much more expensive request mix (a script login costs ~100 ms of a core).

18.1 Response time vs load (closed loop)

Concurrency	Backends	Total	Success	Failed	Err %	Throughput (req/s)	Mean ms	p50	p90	p95	p99	Active	Distribution	Speed up vs 1
1	1	1610	1610	0	0.0	45.99	21.88	11.63	72.37	95.19	134.33	1-1	sys2:100.0%	—
1	2	1538	1538	0	0.0	43.95	22.93	11.98	76.28	97.22	146.66	2-2	sys2:51.0% / sys3:49.0%	0.96x
1	3	1521	1521	0	0.0	43.46	23.16	12.62	75.34	97.06	145.57	3-3	sys2:31.4% / sys3:30.3% / sys4:38.3%	0.94x
10	1	3339	3339	0	0.0	95.4	105.78	68.07	238.59	554.72	840.49	1-1	sys2:100.0%	—
10	2	5736	5736	0	0.0	163.89	61.55	23.88	152.65	267.81	460.72	2-2	sys2:49.0% / sys3:51.0%	1.72x
10	3	7874	7874	0	0.0	224.97	44.75	15.25	122.47	197.3	331.12	3-3	sys2:33.7% / sys3:33.8% / sys4:32.4%	2.36x
25	1	3576	3576	0	0.0	102.16	248.31	89.42	378.24	1704.71	2181.48	1-1	sys2:100.0%	—
25	2	6244	6244	0	0.0	178.42	141.7	69.13	328.54	749.59	1122.3	2-2	sys2:49.9% / sys3:50.1%	1.75x
25	3	9116	9116	0	0.0	260.47	96.99	51.0	253.09	410.63	657.77	3-3	sys2:34.5% / sys3:32.3% / sys4:33.1%	2.55x
50	1	3108	3108	0	0.0	88.81	553.3	80.05	335.24	5186.82	5907.09	1-1	sys2:100.0%	—
50	2	6203	6203	0	0.0	177.23	288.56	105.31	1036.97	1809.18	2366.83	2-2	sys2:48.0% / sys3:52.0%	2.00x
50	3	9618	9618	0	0.0	274.78	183.8	86.37	373.44	982.51	1470.75	3-3	sys2:34.4% / sys3:33.1% / sys4:32.4%	3.09x
100	1	3607	3607	0	0.0	103.06	895.99	110.51	784.04	8976.44	10533.22	1-1	sys2:100.0%	—
100	2	5972	5972	0	0.0	170.63	587.32	168.87	959.22	4402.06	5299.23	2-2	sys2:49.5% / sys3:50.5%	1.66x
100	3	9429	9429	0	0.0	269.4	377.63	126.6	592.07	2587.86	3441.87	3-3	sys2:33.2% / sys3:33.9% / sys4:32.9%	2.61x
200	1	2970	2970	0	0.0	84.85	1384.8	772.25	1356.95	7598.55	12532.98	1-1	sys2:100.0%	—
200	2	6632	6632	0	0.0	189.49	959.87	260.11	1541.22	7855.05	10046.2	2-2	sys2:48.7% / sys3:51.3%	2.23x
200	3	9378	9378	0	0.0	267.93	736.58	216.81	1398.91	5716.81	7009.31	3-3	sys2:32.6% / sys3:34.3% / sys4:33.1%	3.16x




Figure 1 — Response time vs load (p50 left, p95 right), 1 / 2 / 3 backends, log-log.




Figure 2 — Throughput vs load (closed loop).

Reading the table. At one user the three configurations are identical (44–46 req/s, p50 $\approx$ 12 ms): a single sequential client is bound by its own round trip Mac $\rightarrow$ lab $\rightarrow$ Mac, and no number of backends can shorten one request. From ten users upward the backends are the bottleneck and the balancer's value is plain: at 10 users **95 $\rightarrow$ 164 $\rightarrow$ 225 req/s** ($1.7\times$ / $2.4\times$) with p95 falling from 555 ms to 197 ms; at 50 users **89 $\rightarrow$ 177 $\rightarrow$ 275 req/s** ($2.0\times$ / $3.1\times$) with p95 5.2 s $\rightarrow$ 1.0 s; at 200 users **85 $\rightarrow$ 189 $\rightarrow$ 268 req/s** ($2.2\times$ / $3.2\times$) and median response time 772 ms $\rightarrow$ 217 ms. A single 1-core backend saturates at $\sim$ 90–105 req/s regardless of how many users push it — everything beyond that becomes queueing delay (p95 reaches 9 s at 100 users) — while two and three backends move the ceiling to $\sim$ 180 and $\sim$ 270 req/s. The speed-ups slightly above $3\times$ at 50 and 200 users are within run-to-run noise on the shared host (the single-backend runs at those levels happened to land in busier minutes); the honest summary is *near-linear scaling, 2.4–3.2 $\times$ for three backends*. Every run had **zero failed requests**, the adaptive algorithm split traffic 49/51 % and 33/34/33 % (it favours nobody when the backends are equal), and `Active` confirms the balancer saw exactly the intended number of backends throughout each run.

18.2 Throughput vs offered load (open loop)

Offered load nominal → actual (req/s)	Backends	Achieved (req/s)	Success	Failed (of which dropped)	Err %	p50 ms	p95 ms	p99 ms	Active	Distribution
20 → 19.2	1	19.2	480	0 (0)	0.0	23.13	111.11	168.21	1-1	sys2:100.0%
20 → 20.3	2	20.28	507	0 (0)	0.0	23.36	112.24	176.71	2-2	sys2:37.3% / sys3:62.7%
20 → 17.6	3	17.64	441	0 (0)	0.0	27.93	129.19	197.92	3-3	sys2:18.1% / sys3:40.8% / sys4:41.0%
50 → 44.6	1	44.6	1115	0 (0)	0.0	24.83	170.01	339.36	1-1	sys2:100.0%
50 → 43.6	2	43.6	1090	0 (0)	0.0	23.6	135.41	193.85	2-2	sys2:40.8% / sys3:59.2%
50 → 41.2	3	41.16	1029	0 (0)	0.0	23.29	117.69	169.77	3-3	sys2:25.7% / sys3:34.2% / sys4:40.1%
100 → 76.4	1	76.44	1911	0 (0)	0.0	38.99	215.17	768.78	1-1	sys2:100.0%
100 → 79.9	2	79.88	1997	0 (0)	0.0	24.96	159.79	278.7	2-2	sys2:42.1% / sys3:57.9%
100 → 80.5	3	80.48	2012	0 (0)	0.0	25.77	148.35	240.83	3-3	sys2:26.4% / sys3:38.4% / sys4:35.2%
200 → 137.9	1	137.92	3448	0 (0)	0.0	352.22	8948.85	12706.19	1-1	sys2:100.0%
200 → 148.7	2	148.68	3717	0 (0)	0.0	42.08	230.62	417.61	2-2	sys2:48.9% / sys3:51.1%
200 → 145.9	3	145.92	3648	0 (0)	0.0	27.2	181.08	291.68	3-3	sys2:30.5% / sys3:36.5% / sys4:33.1%
300 → 200.2	1	118.64	2966	2038 (797)	40.73	715.03	7422.47	16136.79	1-1	sys2:59.3%
300 → 205.5	2	205.48	5137	0 (0)	0.0	283.48	4488.56	6321.74	2-2	sys2:46.1% / sys3:53.9%
300 → 220.7	3	220.68	5517	0 (0)	0.0	64.08	267.6	415.81	3-3	sys2:33.3% / sys3:33.9% / sys4:32.8%

Actual offered load = arrivals the generator really issued per second (successful + failed + dropped). The thread-per-arrival Python generator on the Mac saturates near 230 arrivals/s, so the nominal 300 req/s setting produced ~220–230 req/s of real offered load.




Figure 3 — Throughput vs offered load (open loop, Poisson arrivals); dashed = ideal.

Reading the table. In open-loop mode arrivals come at a fixed rate whether or not the system keeps up, so the *achieved* curve bends away from the ideal line exactly where the system saturates (Figure 3, left) and

response time explodes just before it (right). Up to ~80 req/s of real offered load all three configurations keep up — throughput follows the ideal line and p95 stays at 110–215 ms. At ~140 req/s the single backend is already over its knee: p50 352 ms and **p95 8.9 s** (a queue that never drains) while two and three backends answer in 231 ms and 181 ms at p95. At ~230 req/s the single backend collapses — **119 req/s achieved, 41 % of arrivals failed** (2 038, of which 797 were shed by the generator's in-flight cap because responses no longer came back), p95 7.4 s — two backends just cope (205 req/s, 0 failed, but p95 4.5 s: at capacity, queue building) and three backends absorb it comfortably (**221 req/s, 0 failed, p95 268 ms**). The offered-load axis uses what the generator actually issued: its thread-per-arrival Python implementation tops out at ~230 arrivals/s, which is why the nominal 300 req/s point is plotted at ~230. Successful/failed request counts per point are in the table; the number of active backends was constant within each run (`Active` column).