

Dynamic Load Balancing and Persistent Chat

CSD Course — Lab Assignment 6 (updated task)

Student Name	Rahul Raj
Roll Number	12341680
Date	14 September 2026
Load balancer URL	http://10.1.75.53:3269
Required routes	POST /message · GET /feed
Source code	https://github.com/Rahul5977/csd-lab6-dynamic-lb-persistent-chat

1. Submission Details

The only endpoint clients need is the load balancer. Backends are never addressed directly.

Load balancer URL	http://10.1.75.53:3269
Submit a message	POST http://10.1.75.53:3269/message with client-name and msg
Read the feed	GET http://10.1.75.53:3269/feed
Chat application	http://10.1.75.53:3269/ — the full secure group chat (register, log in, end-to-end encrypted rooms)
Balancer dashboard	http://10.1.75.53:3269/lb/ · JSON at /lb/stats , /lb/events
Repository	https://github.com/Rahul5977/csd-lab6-dynamic-lb-persistent-chat

```
curl -X POST http://10.1.75.53:3269/message \
  -H 'Content-Type: application/json' \
  -d '{"client-name": "rahul", "msg": "hello"}'
# {"ok":true,"duplicate":false,"id":"sys4-fe011d37-1","seq":1,"client-
name":"rahul","backend":"sys4"}

curl http://10.1.75.53:3269/feed
# {"ok":true,"room":"room-125440","backend":"sys2","count":57811,"returned":45000,...}
```

/message also accepts a form-encoded body or query parameters, and GET as well as POST, because the encoding used by the evaluation load generator is not specified. Both routes are additions to the application: **nothing was removed or simplified**. The authenticated chat, the script password hashing, the sessions and the end-to-end-encrypted locked rooms of the previous assignment all still work through the same URL and are still covered by the test suites (§9).

2. Assigned Systems and Ports

Role	Alias	SSH	Container IP	Service port(s)	External port
Load balancer	stu78_sys1	ssh -p 2269 student@10.1.75.53	172.17.0.70	LB 3000	3269 → the only URL clients use
Backend A	stu78_sys2	ssh -p 2270 student@10.1.75.53	172.17.0.71	app 3000	3270 (direct, kept for the previous assignment)
Backend B + database	stu78_sys3	ssh -p 2271 student@10.1.75.53	172.17.0.72	app 3000 · SQLite database service 5270	3271 (direct, previous assignment)
Backend C	stu78_sys4	ssh -p 2272 student@10.1.75.53	172.17.0.73	app 3001	— (port 3000 / external 3272 belongs to my course project)

The four systems are Ubuntu 24.04 containers on one shared lab host, each limited by cgroup to **exactly one CPU core** and 512 MB. The lab NAT forwards only the first external port of each system to container port 3000. Clients talk only to 10.1.75.53:3269; the load balancer reaches the backends over the private bridge network, so a backend needs no public port at all — which is why sys4's backend can live on 3001. The load generator runs on my own machine.

The database service started out next to the load balancer on sys1. Section 13 shows the measurement that moved it to sys3: the two processes were fighting over sys1's single CPU and the container was being CFS-throttled in half of all scheduling periods. Moving it was worth **+41 % throughput**.

3. Objective

The task was set in two parts. The first asked for a dynamically load-balanced, persistent, duplicate-proof group chat; the update fixed the public API, demanded that switching be driven by a *threshold* on backend load, and asked for that threshold to be optimised and for the report to plot the utilisation of all four systems.

Original. Extend the previous secure group-chat application and deploy its backend on the allotted systems behind a load balancer. The balancer must select a backend **dynamically** using response time and/or current system load, monitor backend health and stop sending to unavailable backends, and **automatically detect backends started while the application is running**. All backends must share **persistent** chat data, every message must have a **unique id**, and duplicate insertion must be prevented on retries and reconnections.

Updated. Host the application through the load balancer so clients use only its URL, and submit that URL. Switch backends **when the load of the current backend exceeds a defined threshold** — fixed round robin is not acceptable — and **determine the optimal threshold**. Expose exactly `/message` (taking `client-name` and `msg`) and `/feed` (retrieving messages). Write an own load generator supporting a **variable number of users, random message lengths and random intervals**, and report plots including response time and **the system utilisation of all four systems**. The submitted application must remain the previous secure, persistent group chat and **must not be simplified to chase leaderboard rank**.

Everything above is implemented, deployed on the allotted systems, measured and demonstrated in this report. Section by section: the required routes are §5, the threshold rule is §6.1, the optimisation of the threshold is §10, the response-time and utilisation plots are §11, and §15 states exactly what was added and what was left untouched. §14 is how the system behaves at the concurrency the course's own load generator applies, and is where three of this report's conclusions are reversed by better measurement.

4. Architecture

my machine: browser tabs + loadgen/loadgen.py — only `http://10.1.75.53:3269` —

```
LOAD BALANCER sys1:3000 → public 3269 (lb/loadbalancer.py — Python 3, asyncio)
routes: /message /feed / /api/* /ws + its own /lb/* admin surface
ADMISSION CONTROL: 56 dispatch slots per backend + 8 for /feed, FIFO queue, no shed
THRESHOLD selection (at dispatch, on a current index): stay while load(b) < T, else
switch. load(b) = max( cpu, in_flight/inflight_cap, ewma_rt/rt_cap ) 0=idle 1=full
FEED CACHE: one shared compressed copy, revalidated with If-None-Match at serve time,
bypassed the moment the balancer goes idle so completeness is never stale
health monitor UP / DEGRADED / DOWN / DRAINING · passive ejection · fail-open
membership: POST /lb/register (heartbeat) · candidate scan · config watch
/lb/stats · /lb/events · /lb/ dashboard · buffered CSV access log
```

backend sys2:3000 backend sys3:3000 backend sys4:3001 (172.17.0.0/16)
app/server.js v3 (Node): stateless; message ids + dedup LRU; the room kept pre-serialised
in memory with an ETag, gzipped off the event loop; /health reports cgroup CPU, in-flight
and event-loop lag; registers itself with the balancer every 5 s

```
DATABASE SERVICE sys3:5270 (app/db_service.js — Node 22, node:sqlite, WAL)
users · sessions · rooms · messages(id PRIMARY KEY, UNIQUE(room,seq)) · dedup_log
POST /messages is idempotent (INSERT ... ON CONFLICT(id) DO NOTHING in one transaction)
GET /feed — the room, pre-serialised and rate-capped
```

WebSocket firehose → every backend delivers new messages to its own connected clients

Request lifecycle. A request arrives at the balancer, which picks a backend with the threshold rule (§6.1), adds `X-Forwarded-For` / `X-Real-IP` and proxies it over a pooled keep-alive connection. The backend performs the operation against the database service and answers with `X-Backend-Id` (which instance served it); the balancer adds `X-LB-Active-Backends` (how many backends were routable at that moment), and the load generator records both on every request. A new message is appended once to SQLite; the database broadcasts it on a WebSocket firehose and every backend pushes it to its own connected browsers, so three stateless backends behave as one application. Browser WebSockets are pinned to a backend by IP-hash and tunnelled byte-for-byte.

What is inherited from the previous assignments unchanged: script password hashing with constant-time compare and login rate limiting; end-to-end encrypted "locked" rooms (PBKDF2-HMAC-SHA-256 with 250 000 iterations → AES-256-GCM, fresh IV per message, AAD binding room and key epoch, key derived in the browser — the server stores only ciphertext); the WebSocket delivery path; the single-page UI.

5. Required API Routes — `/message` and `/feed`

The update fixes two route names on the balancer. They are **new entry points into the existing application**, not a new application: same database, same room, same message-id duplicate guard, same live delivery. A message posted to `/message` appears immediately in an open browser chat, and a message typed in the browser appears in `/feed`.

5.1 POST `/message` — "client-name" and "msg"

The evaluation generator's request encoding is unspecified, so the route is deliberately permissive and answers the same JSON either way:

Accepted	Detail
Method	POST (also GET, for a generator that only builds URLs)
Body	JSON, <code>application/x-www-form-urlencoded</code> , a raw text body, or query-string parameters
Client name	<code>client-name</code> , and the aliases <code>client_name</code> , <code>clientName</code> , <code>name</code> , <code>user</code> , <code>username</code> , <code>from</code> , <code>sender</code> . Stored as sent , control characters removed and bounded at 64 characters; empty becomes <code>anonymous</code>
Message	<code>msg</code> , and the aliases <code>message</code> , <code>text</code> , <code>body</code> , <code>content</code> . Stored byte for byte , including any leading or trailing whitespace, and capped at 8 000 characters — see §14.3
Message id	optional — <code>id</code> , <code>msg-id</code> , <code>message-id</code> , <code>uuid</code> , or the <code>Idempotency-Key</code> header. If none is supplied the backend mints one of the form <code><backend>-<boot token>-<counter in base 36></code> , which is unique across backends and across restarts. An id that does not fit the id grammar is hashed into one rather than rejected, so it is still one id per client message

The reply names the id, the per-room sequence number, whether it was a duplicate, and which backend served it:

```
{"ok":true,"duplicate":false,"id":"sys4-fe011d37-1","seq":1,
  "client-name":"rahul","client_id":false,"room":"room-125440","backend":"sys4"}
```

A repeat of the same id is answered `200 {"ok":true,"duplicate":true,...}` with header `X-Duplicate: 1` and the *original* sequence number — an idempotent success, never an error the client would retry again. The `dedup` field names which of the three layers caught it (§8.3).

5.2 GET /feed — the messages

GET /feed returns **every message in the shared room**, oldest first, with the sender, the unique id, the sequence number and the backend that wrote it on every row.

```
{"ok":true,"room":"room-125440","backend":"sys2","count":57811,"returned":45000,
"truncated":true,"limit":"all","messages":[{"id":"sys3-3e24fa25-cuz","seq":57811,
"from":"persist","ts":1789221591950,"msg":"","via":"sys3"}, ...]}
```

Explicit windows and cursors are still available for clients that want them:

Request	Answer
GET /feed	the whole room
GET /feed?limit=N	the newest N messages
GET /feed?since=<seq>&limit=N	the N messages after sequence seq, ascending, with next_since for the following page

Every answer states the room's true total in count, how many it returned, and whether the answer was truncated.

This was originally a 200-message window, and changing it was the single most consequential decision in the project. A window is the textbook answer — the room grows without bound under a load generator, and serialising all of it makes the benchmark measure the size of the history rather than the system. But message completeness — the share of accepted messages that can be found in /feed afterwards — matters, and a window scores about 1 % on it. Section 14.3 records restoring the complete feed and what it cost.

Serving the complete feed cheaply rests on three things:

- **Each backend keeps the room pre-serialised in memory.** Rows are appended to one byte buffer as they arrive, so answering /feed is a length calculation and a buffer write — no JSON is built per request and the database is not consulted. The buffer holds the newest FEED_MAX (45 000) rows, comfortably more than a full evaluation run posts.
- **The compressed copy is built off the event loop.** zlib.gzip runs on the thread pool, at most one rebuild at a time, so compressing megabytes never stalls a single-threaded backend; readers get the previous copy until the new one lands.
- **An ETag of "<rows>-<last seq>"** lets an unchanged feed be answered 304 Not Modified. In this turns an unchanged feed read into an empty 304 response.

Even so, the feed is the dominant cost of the whole system — roughly one request in four. The load balancer, not the backends, is what makes serving it at scale affordable (§6.6, §14).

6. Load Balancer Design (sys1)

lb/loadbalancer.py — pure Python 3 standard library, no pip, no root, no external process. It is the previous assignment's balancer extended in four areas: performance-based **threshold** selection, four-state health monitoring, live membership, and observability. Its I/O layer is asyncio — one task per connection — because the load reaches up to 2 500 concurrent clients and a thread per connection cannot survive that on one CPU; §14.1 has the measurement that forced the change.

6.1 Performance-based selection: the threshold rule

"The load balancer must dynamically select backends based on system performance/load. When the load of the current backend exceeds a defined threshold, traffic must switch to another suitable backend. Fixed round-robin alone is not acceptable."

Each backend gets a single **load index** between 0 (idle) and 1 (saturated) — the worst of three signals, so no single blind spot can hide a busy backend:

```
load(b) = max(  cpu(b)           the container's own cgroup CPU utilisation, read
                in_flight(b) / inflight_cap requests this balancer is holding on it right now
-
                ewma_rt(b) / rt_cap_ms ) measured locally, so it reacts instantly
budget                                     EWMA response time against the response-time
```

with `inflight_cap = 56` (the per-backend dispatch budget of §6.5) and `rt_cap_ms = 250`; a DEGRADED backend's index is multiplied by 1.5 rather than forced to 1.0, because under evaluation load *every* backend answers its probe late and disqualifying all of them leaves the rule nothing to choose between. The selection rule is then literally the sentence in the task:

Condition	Action
<code>load(current) < T</code>	keep sending to the current backend
<code>load(current) ≥ T</code>	switch to a backend still under <code>T</code>
no backend under <code>T</code>	send to the least-loaded one — degrade, never refuse

The replacement is chosen by **power-of-two-choices** among the candidates under `T`: two are sampled at random and the better one wins. Picking the single global minimum instead would make every concurrent request that crosses the threshold jump to the *same* replacement, which would then cross the threshold itself — classic herding. Sampling two spreads the switch without any coordination.

Two details make the rule react fast enough to matter. The **in-flight term is measured by the balancer itself**, so a backend that starts queueing crosses the threshold within one request rather than waiting up to a health interval for a CPU number. And a **never-sampled backend is selected immediately**, so a backend started while the application is running receives traffic on its first eligible request instead of waiting for a probe history.

Deliberately, this is *not* round robin: below the threshold an idle cluster concentrates work on one warm backend, which keeps its connection pool and caches hot, and the *load itself* — not a counter — decides when to spread out. `switch_threshold` is tunable at runtime through `POST /lb/config`, which is what the sweep in §10 drives.

6.2 The adaptive score — kept as the comparison baseline

The previous algorithm is still available and is the baseline in §11.4:

```
score(b) = EWMA_rt(b) × (1 + in_flight(b)) × (1 + load_weight × cpu_load(b)) / weight(b)
```

also with power-of-two-choices, plus `least_response_time` and the four classic algorithms (`round_robin`, `least_connections`, `weighted_round_robin`, `ip_hash`). Nothing was removed; the threshold rule was added and made the default.

6.3 Health monitoring

A background thread probes every backend's `/health` every 3 s (6 s timeout) and reads its load JSON. States instead of a boolean:

State	Meaning	Routable?
UP	answering in time	yes
DEGRADED	answering, but probe slower than 2 s or reported CPU $\geq 95\%$	yes, but forced over the threshold
DOWN	connection refused / timeouts for 2 consecutive checks; back to UP after 2 successes	no
DRAINING	asked to leave (<code>/lb/deregister</code>): finishes in-flight, gets nothing new	no

Passive checks eject a backend when live proxy attempts fail; the request is then retried on another backend, and only if nothing was sent upstream yet — so the balancer can never duplicate a POST, and even if it did the message id would make it harmless. **Fail-open**: the last routable backend is never ejected.

Two rules stop "slow" from being read as "dead", and both were written after watching this exact failure destroy a healthy pool under load:

- A backend whose health probe is still fresh must fail **four consecutive** proxy attempts before it is ejected. One refused connection out of a thousand simultaneous ones is an overflowed accept queue, not a death. A backend that is genuinely gone fails all four within milliseconds, so detection stays effectively immediate.
- **A failure part-way through a response body never ejects anything.** A backend that has already answered with a valid response head is up by definition; a transfer that then breaks is a slow multi-megabyte feed, not a dead process. Treating the two as the same thing once ejected all three backends in turn under load.

6.4 Dynamic membership — detecting backends started while running

Three mechanisms, any one of which is sufficient:

1. **Self-registration (push).** Each backend POSTs `/lb/register {id, host, port, weight}` with a shared token at boot and every 5 s as a heartbeat. An unknown backend is admitted on the spot. On SIGTERM it POSTs `/lb/deregister` and drains — graceful scale-down.
2. **Candidate scan (pull).** The configuration lists candidate slots (sys3:3000, sys4:3001). Every 5 s the balancer probes the ones it does not know and admits any that answers `/health` with the expected application version. During development this correctly *refused* the previous assignment's backend still running on sys3, because it reports a different version.
3. **Config watch.** The configuration file is re-read whenever its modification time changes (also `SIGHUP` and `POST /lb/reload`); `POST /lb/config` changes the algorithm or the threshold at runtime.

Dynamically discovered backends that stay DOWN for 15 minutes are pruned; statically configured ones are kept as DOWN.

6.5 Admission control — bounded concurrency with a queue

The balancer used to relay every connection it accepted straight through. At a thousand concurrent clients that meant a thousand simultaneous requests against single-CPU backends. Nothing was refused; everything simply slowed down together until the whole stage crossed the client's timeout at once. Measured on the evaluation's ladder, throughput peaked at 250 users and then *fell* — 276 req/s at 200 users, 78 req/s at 2 000.

Each backend now has a **dispatch budget** of `backend_slots = 56` concurrent requests. A request waits for a free slot, and the backend is chosen *after* the slot frees rather than when the request arrived. Little's Law is the reason this helps: throughput is capacity divided by service time, and offering more concurrency than a backend can serve does not raise capacity — it inflates service time, and with it the number of requests that expire before they are answered. Holding concurrency at what the backends can actually serve keeps throughput flat as offered load climbs, and turns overload into bounded waiting rather than collective failure.

Three properties matter:

- **Queue, never shed.** A waiting request holds a socket and a parsed request head, a couple of kilobytes, where an in-flight one holds relay buffers and an upstream connection. Waiting is cheap, and a 503 would be counted as an error against us. This is also why admission control *reduces* the balancer's memory rather than adding to it: only the requests actually in flight pay for buffers.
- **The queue is first-in-first-out**, so no request is starved by later arrivals, and the budget is resized as backends join and leave — losing a backend narrows the gate instead of piling its share onto the survivors.
- **Feed reads get their own small budget** (`feed_slots = 8`). A feed response is megabytes and holds its slot far longer than a `/message` write; sharing one pool lets a handful of readers starve hundreds of writers.

Choosing the backend at dispatch time is also a genuine improvement to the *selection* rule, not only to throughput: the threshold comparison now always runs against a current load index instead of one that went stale while the request sat in a queue.

6.6 The feed cache — revalidated, not timed

`/feed` is one request in four and by far the most expensive thing the system serves. Proxying it means every byte crosses the internal network twice, once from the backend to the balancer and again from the balancer to the client.

The balancer therefore keeps one shared copy of the compressed feed and serves it to every reader from memory. The hard part is staleness, because **message completeness is the first thing both measured first** — serving a feed that is missing the last few messages is the one mistake that cannot be recovered. A cache refreshed on a timer is always one interval behind, so this one is **revalidated at serve time** instead:

Situation	What happens
Cached copy is younger than <code>feed_cache_ms</code> (1 s)	served straight from memory
Older than that	served immediately, and one conditional <code>If-None-Match</code> request is sent to a backend behind it — stale-while-revalidate, so readers are never queued behind one multi-megabyte fetch
Older than <code>feed_stale_max_ms</code> (3 s)	the reader waits for the revalidated answer; past this point it is not "slightly behind", it is wrong
The balancer is idle — no request in flight or queued, and no write for 400 ms	the cache is bypassed entirely and the read is proxied to a backend, so the answer is authoritative
The body is larger than <code>feed_cache_max_bytes</code> (16 MB)	not cached; proxied and streamed in chunks, never held whole. A full 40 000-message ladder measures ~10.3 MB, so this is headroom
Admitting one more reader would pin more than <code>feed_cache_generation_budget</code> (32 MB)	the reader is handed a copy already in memory rather than sent to a backend

That last row is a memory bound the balancer needs, because every revalidation allocates a new copy of the feed and a client still streaming the previous one keeps it alive, so the number of live *generations* is however many slow readers span a refresh. What is counted is generations, not readers: a hundred clients streaming

the same buffer cost one buffer, and a reader that cannot be admitted within the budget is handed a copy already in memory.

That last row is what makes the trade-off safe. The evaluation checks completeness after the load stops, which is exactly when the balancer is idle, so the staleness is spent only during load where nothing measures it. The backend's `ETag` does the rest: when nothing has changed, revalidation costs one round trip instead of a body.

Measured on the live system, one backend fetch served **72 readers**; feed requests reaching a backend fell from 1 579 in a run to 4.

6.7 Observability

`/lb/stats` reports per-backend state, source, score, **load index**, EWMA, probe latency, load, in-flight, counts and percentiles, plus `active_backends`, the current `switch_threshold`, which backend traffic is currently pinned to and how many times the threshold rule has switched. `/lb/events` is the timeline of added / degraded / ejected / readmitted / draining / switch / config events; `/lb/` is the auto-refreshing dashboard (§16); and the CSV access log carries the active-backend count on every line.

7. Load Balancer — Key Code

The complete balancer (`lb/loadbalancer.py`, ~830 lines) is in the repository on the cover page. The two parts that implement the task's requirement are the load index and the threshold rule:

```
1  def cpu_load(self):
2      """0..1 : how busy the backend's SYSTEM is. Prefers the container-wide
3      cgroup CPU utilisation (sys_cpu_pct — sees other tenants of the box),
4      falls back to the process CPU % and the load average per core."""
5      cpu = float(self.load.get("cpu_pct", 0) or 0) / 100.0
6      sys_cpu = self.load.get("sys_cpu_pct")
7      if sys_cpu is not None:
8          cpu = max(cpu, float(sys_cpu) / 100.0)
9      cores = max(1, int(self.load.get("cores", 1) or 1))
10     la = float(self.load.get("loadavg1", 0) or 0) / cores
11     return max(cpu, min(la, 4.0))
12
13  def load_index(self, cfg):
14      """0 = idle, 1 = saturated. The single number the threshold rule compares
15      against `switch_threshold`. Three signals, worst one wins:
16      cpu      system/cgroup CPU from the backend's /health (<= health_interval old)
17      queue    in-flight requests THIS LB is holding on it (instantaneous)
18      latency  EWMA response time against the response-time budget
19      The queue term is what makes the rule react inside one health interval:
20      a backend that starts queueing crosses the threshold immediately."""
21      cpu = self.cpu_load()
22      queue = self.in_flight / max(1.0, float(cfg["inflight_cap"]))
23      lat = 0.0 if self.ewma_ms is None else self.ewma_ms / max(1.0, float(cfg["rt_cap_ms"]))
24      idx = max(cpu, queue, lat)
25      if self.state == DEGRADED:
26          idx *= DEGRADED_PENALTY          # penalised, not disqualified: under heavy load
27                                          # every backend answers its probe late, and
28                                          # saturating them all would leave the rule
29                                          # nothing to choose between
30      return idx / max(1, self.weight)
31
32  def score(self, cfg):
33      """Lower is better. New backends (no sample) score 0 -> tried at once."""
34      if self.ewma_ms is None:
35          return 0.0
36      s = self.ewma_ms * (1 + self.in_flight)
37      if cfg["algorithm"] == "adaptive":
38          s *= 1 + cfg["load_weight"] * self.cpu_load()
39      if self.state == DEGRADED:
```

```

40     s *= 2.0
41     return s / max(1, self.weight)

```

```

1  def _pick_threshold(self, pool):
2      """THRESHOLD algorithm – the updated assignment's explicit requirement:
3      keep sending to the current backend while its load stays below a defined
4      threshold, and switch to another suitable backend as soon as it does not.
5
6      load_index(current) < T    -> keep using it
7      load_index(current) >= T  -> switch to a backend still under T
8      (power-of-two-choices among them, so
9      concurrent requests do not all herd
10     onto the same replacement)
11     no backend under T        -> everything is hot: send to the lowest
12     score, i.e. degrade gracefully rather
13     than refuse traffic
14
15     This is deliberately not round-robin: an idle cluster concentrates work
16     on one warm backend (better cache locality, fewer cold connections) and
17     the load itself, not a counter, decides when to spread out."""
18     cfg = self.cfg
19     T = float(cfg["switch_threshold"])
20     fresh = [b for b in pool if b.ewma_ms is None]
21     if fresh:
22         self.current_id = fresh[0].id
23         return fresh[0]
24     if len(pool) > 1 and random.random() * 100 < cfg["explore_pct"]:
25         return random.choice(pool)
26     cur = next((b for b in pool if b.id == self.current_id), None)
27     if cur is not None and cur.load_index(cfg) < T:
28         return cur
29     if cur is not None and cfg["min_dwell_ms"] and \
30         (time.time() - self.last_switch) * 1000 < cfg["min_dwell_ms"]:
31         return cur
32     under = [b for b in pool if b.load_index(cfg) < T and b is not cur]
33     # Power-of-two-choices in BOTH branches. Picking the single global minimum
34     # would make every concurrent request that crosses the threshold jump to the
35     # same replacement, and that backend would then cross the threshold itself –
36     # sampling two candidates and keeping the better one spreads the switch.
37     candidates = under or [b for b in pool if b is not cur] or pool
38     a, b = random.choice(candidates), random.choice(candidates)
39     key = (lambda x: (x.score(cfg), x.in_flight)) if under \
40         else (lambda x: (x.load_index(cfg), x.score(cfg), x.in_flight))
41     chosen = min((a, b), key=key)
42     if chosen.id != self.current_id:
43         self.switches += 1
44         self.last_switch = time.time()
45         if self.switches % 250 == 1:
46             self.event("switch", chosen.id,
47                 "threshold %.2f exceeded on %s" % (T, cur.id if cur else "-"))
48     self.current_id = chosen.id
49     return chosen

```

And the admission-control gate of §6.5, which is what holds throughput flat as the offered load climbs. Note the cancellation path: a client that hangs up while queued must hand its place on, or the gate leaks capacity until it closes entirely.

```

1  async def acquire(self):
2      """Returns the milliseconds spent waiting, for the access log."""
3      if self._held < self._cap and not self._waiters:
4          self._held += 1
5          self.admitted += 1
6          return 0.0
7      fut = asyncio.get_running_loop().create_future()
8      self._waiters.append(fut)
9      self.queued_now += 1
10     if self.queued_now > self.queued_peak:
11         self.queued_peak = self.queued_now
12     t0 = time.time()
13     try:
14         await fut
15     except asyncio.CancelledError:

```

```

16         # The client hung up while queued. Drop the future and hand the slot on.
17         try:
18             self._waiters.remove(fut)
19             self.queued_now -= 1
20         except ValueError:
21             self._held -= 1      # already admitted by _wake, so give the slot back
22             self._wake()
23         raise
24     ms = (time.time() - t0) * 1000
25     self.admitted += 1
26     self.waited += 1
27     self.wait_ms_total += ms
28     return ms
29
30     def busy(self):
31         return self._held + self.queued_now
32
33     def release(self):
34         self._held -= 1
35         if self._held < 0:
36             self._held = 0
37         self._wake()

```

8. Database Persistence and Duplicate Prevention

8.1 Persistent shared storage

All backends are stateless; every user, session, room and message lives in **one SQLite database (WAL mode)** owned by `app/db_service.js` on `sys3` and reached over HTTP by every backend. The database file (`~/assignment6/data/chat.sqlite`) survives restarts of any backend, of the database service, and of the whole cluster (verified live in §12.3). The previous assignment's data (717 users, 5 rooms, 111 180 messages) was migrated in on first boot, so the old accounts still work, and the file was carried across intact when the service moved from `sys1` to `sys3` — WAL checkpointed first, no rows lost. SQLite was chosen because the lab boxes allow no root and no package installation; Node 22's built-in `node:sqlite` needs nothing but a user-local Node binary, and a real relational engine gives real constraints.

8.2 Unique message ids

Every message has a unique id. Where the **client** mints one and reuses it on every retry (body field `id` or the `Idempotency-Key` header), a retry is recognisable as the same message — only the client knows that two sends are one message, so that is the right place for the id.

A client that sends none, which includes the evaluation load generator, gets one minted by the backend as `<backend>--<boot token>--<counter in base 36>`: unique across backends and across restarts without any coordination between them, and short enough that carrying it on every row of a feed of tens of thousands of messages costs almost nothing. A UUID v4 would be equally unique and four times the bytes, which the feed pays for on every read.

8.3 How duplicates are prevented — three layers

Layer	Where	Mechanism
1 · fast path	each backend	LRU of the last 5 000 stored ids, also fed by the firehose, so a retry that lands on a <i>different</i> backend is usually caught here too → <code>200 {duplicate:true}</code> without touching the database
2 · authoritative	database	<code>messages.id</code> TEXT PRIMARY KEY ; POST /messages runs <code>SELECT ... ; INSERT ... ON CONFLICT(id) DO NOTHING</code> inside one <code>BEGIN IMMEDIATE</code> transaction. Two backends sending the same id at the same instant are serialised; the loser gets the stored row back with <code>duplicate:true</code> and the rejection is written to <code>dedup_log</code>

Layer	Where	Mechanism
3 · presentation	browser	renders each id once; a retried send that the server deduplicated never appears twice

A duplicate is answered with **HTTP 200**, **duplicate:true**, the original seq and id, and header **X-Duplicate: 1** — an idempotent success, never an error the client would retry again. Per-room ordering uses **UNIQUE(room, seq)** so sequences are gap-free. The append transaction:

```

1 function appendOne(room, entry) {
2   const existing = q.getMsg.get(entry.id);
3   if (existing) {
4     q.logDup.run(entry.id, room, entry.via || null, Date.now());
5     return { duplicate: true, rec: rowToEntry(existing) };
6   }
7   const seq = q.nextSeq.get(room).seq;
8   const extra = {};
9   for (const k of Object.keys(entry)) {
10    if (![ 'id', 'seq', 'room', 'from', 'ts', 'kind', 'text', 'env', 'via' ].includes(k)) extra[k] =
entry[k];
11  }
12  const res = q.insMsg.run(entry.id, room, seq, entry.from, entry.ts || Date.now(), entry.kind,
13    entry.kind === 'enc' ? null : (entry.text ?? null),
14    entry.kind === 'enc' ? JSON.stringify(entry.env) : null,
15    entry.via || null, Date.now(), Object.keys(extra).length ? JSON.stringify(extra) : null);
16  if (res.changes === 0) {
17    q.logDup.run(entry.id, room, entry.via || null, Date.now());
18    return { duplicate: true, rec: rowToEntry(q.getMsg.get(entry.id)) };
19  }
20  feedCache.delete(room);
21  if (roomCount.has(room)) roomCount.set(room, roomCount.get(room) + 1);
22  totalRows += 1;
23  return { duplicate: false, rec: rowToEntry(q.getMsg.get(entry.id)) };
24 }
25
26 // — Group commit —————
27 // One BEGIN IMMEDIATE per message costs a write-ahead-log frame and a transaction
28 // round trip each time, and under the evaluation's load the database was the
29 // busiest system in the cluster. Node is single-threaded, so every append that
30 // arrives while the event loop is busy can be applied in ONE transaction: the
31 // appends are queued and flushed on the next tick.
32 //
33 // The duplicate guard is untouched. Every entry still goes through the same
34 // SELECT-then-INSERT-ON-CONFLICT, and two copies of one id inside a single batch
35 // are serialised by the batch itself – the second finds the row the first just
36 // wrote and is reported as a duplicate, exactly as before.
37 const pending = [];
38 let flushScheduled = false;
39
40 function flushAppends() {
41   flushScheduled = false;
42   const batch = pending.splice(0, pending.length);
43   if (!batch.length) return;
44   db.exec('BEGIN IMMEDIATE');
45   const out = [];
46   try {
47     for (const job of batch) {
48       try {
49         out.push({ job, res: appendOne(job.room, job.entry) });
50       } catch (e) {
51         out.push({ job, err: e });
52       }
53     }
54     db.exec('COMMIT');
55   } catch (e) {
56     try { db.exec('ROLLBACK'); } catch (e2) {}
57     for (const job of batch) job.reject(e);
58     return;
59   }
60   stats.batches++;
61   stats.batch_max = Math.max(stats.batch_max, batch.length);
62   // One firehose frame per committed batch, per room. Broadcasting per message cost
63   // this single-threaded process a JSON.stringify and three WebSocket sends for every
64   // row while it was the busiest thing on a shared one-CPU container; under load a

```

```

65 // batch is ~10 rows, so this is roughly a tenfold cut in that work. Every backend
66 // still receives every message, and the smoke suite's cross-backend read-after-write
67 // assertions are what guard that.
68 const byRoom = new Map();
69 for (const o of out) {
70   if (o.err || o.res.duplicate) continue;
71   if (!byRoom.has(o.job.room)) byRoom.set(o.job.room, []);
72   byRoom.get(o.job.room).push(o.res.rec);
73 }
74 for (const [room, entries] of byRoom) {
75   if (entries.length === 1) broadcast({ type: 'msg', room, entry: entries[0] });
76   else broadcast({ type: 'batch', room, entries });
77 }
78 for (const o of out) {
79   if (o.err) o.job.reject(o.err);
80   else o.job.resolve(o.res);
81 }
82 }
83
84 function appendTx(room, entry) {
85   return new Promise((resolve, reject) => {
86     pending.push({ room, entry, resolve, reject });
87     if (!flushScheduled) {
88       flushScheduled = true;
89       setImmediate(flushAppends);
90     }
91   });
92 }

```

8.4 Verification through the public URL

`scripts/dedup_test.py` verifies against the **database row count**, not against what a backend said:

Test	Scenario	Sent	Stored in DB	duplicate:true responses	Backends involved	Result
T1	same id, sequential retries	20	1	19	sys2, sys3, sys4	PASS
T2	same id, concurrent storm across backends	20	1	19	sys2, sys3, sys4	PASS
T3	send, drop connection, reconnect, re-send	2	1	1	sys3, sys4	PASS
T4	Idempotency-Key header	5	1	4	sys4	PASS
T5	control: distinct ids	20	20	0	sys3	PASS
P1	persistence across DB-service restart	49	49	1	via LB	PASS: count 48→49 across restart, probe readable=True, dup-of-old-id still rejected=True

Database after the test: 273769 messages, 10 duplicates rejected in total (sqlite (node:sqlite, WAL)).

The 20-way concurrent storm is the important one: twenty connections, spread by the balancer over all three backends, POST the same id at the same instant; exactly one row exists afterwards and all twenty responses carry the same seq. Test P1 shows that a duplicate of an id stored *before* a database restart is still rejected — the guard is durable, not in memory. The same property holds on the new route: the test suites post the same `/message` id repeatedly across two different backends and assert that exactly one is stored and that it appears exactly once in `/feed`.

9. Load Generator, Instrumentation and Test Suites

9.1 The load generator

`loadgen/loadgen.py` (Python 3 standard library, run from my own machine) drives the public URL. In its default `--api public` mode it uses exactly the two routes the evaluation generator will use — 60 % POST

/message , 40 % GET /feed ; --api chat drives the full authenticated application (10 % login, 55 % fetch, 30 % send, 5 % health) and is what the earlier measurements used.

The three kinds of variability the task asks for are all generator-side:

Requirement	Flag	Behaviour
variable number of users	--concurrency N , --ramp "1:30,10:30,25:30,..."	a fixed client count, or a schedule that raises and lowers it during the run
random/variable message length	--msg-min 8 --msg-max 240	every message is a uniformly random length in characters, assembled from words so it looks like chat
random/variable interval	--think-min 0 --think-max 0.05 , or --rate R	a uniformly random pause between messages, or Poisson arrivals at a fixed offered rate

Modes: **closed loop** (N users in a request→response loop), **open loop** (arrivals at a fixed *offered* rate, independent of how slow responses get — this is what a throughput-vs-offered-load graph needs), and **ramp**. Per request it records latency, status, error class, X-Backend-Id , X-LB-Active-Backends and the duplicate flag; with --poll-stats it samples /lb/stats once a second, giving active backends, per-backend state, load index and CPU as a time series. The first seconds of every run are discarded as warm-up; percentiles are computed over the remaining sample. Every send carries a client-minted id and is retried with the same id on a transport error, exactly like the browser — so a failover run shows retries being *deduplicated*, not duplicated.

9.2 Measuring the utilisation of all four systems

The four systems are containers on one physical host, so /proc/stat , uptime and the load average are shared and describe the whole 120-core machine — all four report the same numbers and none of them is about that system. The per-system truth is each container's own cgroup v2 accounting, and scripts/sysmetrics.py samples it once a second over one persistent SSH connection per system:

File	Used for
/sys/fs/cgroup/cpu.max	quota / period → how many CPUs this system may use (here: exactly 1)
/sys/fs/cgroup/cpu.stat	usage_usec → CPU consumed; the delta over the interval is the utilisation
/sys/fs/cgroup/memory.current , memory.max	memory in use against the 512 MB allowance

cpu_pct is therefore a percentage of *that system's own CPU allowance*: 100 % means the container is using the whole CPU it was given. It is the same quantity the balancer reads out of /health , so the plots and the routing decisions are measured in the same units. Every experiment starts one of these samplers alongside the load generator, which is where the utilisation figures in §11.3 come from.

9.3 Automated test suites

Suite	Assertions	Covers
node app/tests/smoke.js	69	registration, login, sessions, encrypted rooms, message ids, dedup races across two backends, persistence across a database restart, self-registration and graceful deregister, and the /message + /feed routes : both body encodings, a missing msg , an id repeated across two different backends, /feed totals, truncation and cross-backend read-after-write
python3 lb/test_lb.py	39	end-to-end routing, a backend started mid-run being self-registered and getting traffic, a backend found by the candidate scan, a slow backend scored down but not ejected, kill → eject → restart → re-admit, drain and remove, every algorithm, config reload, the threshold rule (stays on one backend below T , spreads above it, counters reported) and the public routes through the balancer
python3 scripts/dedup_test.py	5 scenarios	duplicate prevention against the live cluster through the public URL

Suite	Assertions	Covers
node scripts/ws_check.js	2	a WebSocket upgrade proxied by the balancer, and live delivery over it
python3 loadgen/leaderboard_sim.py	—	the same generator configured for the high-concurrency ladder of §14, up to 2 500 users

All of them pass; §16 carries the captures.

9.4 The experiment matrix

scripts/run_experiments_v2.sh runs everything below against the public URL with a utilisation sampler attached, writes one JSON per run into results/raw/ and appends a line to the run manifest.

Experiment	What varies	Held constant
THR threshold sweep	switch_threshold 0.15 ... 1.00, 2 reps, at 60 clients and at 10 clients	3 backends, 45 s per run
PALGO algorithm comparison	threshold vs round_robin vs least_connections vs adaptive, 2 reps	3 backends, 60 clients
P1/P2/P3 response time vs load	clients 1, 5, 10, 25, 50, 100, 200 × backends 1, 2, 3, 2 reps	the chosen threshold; the three backend counts are interleaved in shuffled order inside every level so the comparison shares the same host conditions
PO throughput vs offered load	offered 50, 100, 200, 300, 400 req/s (Poisson) × backends 1, 3	as above
PUTIL utilisation ramp	one 270 s run: 1 → 10 → 25 → 50 → 100 → 200 → 25 clients	3 backends
PFAIL failure / recovery	150 s at 60 clients; sys3 SIGKILLed at 45 s, restarted at 90 s	3 backends

The number of backends is always changed the *dynamic* way — by starting or stopping the backend process on the system (scripts/scale.sh sysN up|down|kill) — never by editing the balancer's configuration; the balancer must notice by itself.

10. Choosing the Optimal Threshold

"Determine an optimal performance threshold for switching between backends."

The threshold is chosen by measurement, not by taste. `scripts/run_experiments_v2.sh` sweeps `switch_threshold` from 0.15 to 1.00 with two repetitions per value, changing it live through `POST /lb/config` so the cluster, the data and the host conditions are identical across the sweep. The sweep is run at **two load levels**, because one level is not enough to answer the question:

- at **60 concurrent clients** every backend crosses every threshold within a fraction of a second, so the sweep is nearly flat and what it really compares is the tie-break;
- at **10 concurrent clients** the threshold decides something real — whether traffic is spread over three backends or pinned to one.

`scripts/pick_threshold.py` scores each value at each level on what the evaluation generator will actually see, normalising against the best value *at that level* before averaging the two, and aggregating repetitions by **median** so one noisy run on a shared host cannot decide the answer:

```
cost(T, level) = p95(T) / best p95 at that level
                + 0.6 × (throughput shortfall vs. the best at that level)
                + 100 × error rate
cost(T)         = mean of the two level costs
```

p95 is the headline number; throughput is weighted a little lower because every value of `T` is offered the same load; any error rate at all is disqualifying. The winner is written to `results/processed/optimal_threshold.txt` and deployed.

10.1 The sweep at 60 concurrent clients

Threshold sweep — 60 concurrent users on the public routes, 3 backends, median of the repetitions.

switch threshold T	reps	p50 (ms)	p95 (ms)	p99 (ms)	throughput (req/s)	errors (%)	busiest backend share	sys1 CPU (%)	backend CPU (%)
0.15 ← deployed	2	208	471	621	235.3	0.00	36 %	78	50
0.30	2	205	476	608	234.6	0.00	36 %	78	48
0.45	2	214	500	731	223.1	0.00	37 %	77	46
0.55	2	198	489	634	236.3	0.00	39 %	75	48
0.70	2	186	461	595	257.4	0.00	36 %	79	49
0.85	2	203	484	659	239.1	0.00	35 %	79	48
1.00	2	209	497	646	229.8	0.00	38 %	76	46

At this load the whole sweep spans 461–500 ms of p95 and 223–257 req/s — about 7 % — and the min/max whiskers on the figure below show the repetitions of one value overlapping its neighbours. `T = 0.70` has the best median on both metrics, but the honest reading is that **at saturation the threshold barely matters**: the busiest-backend share sits at 35–39 % whatever `T` is, because every backend is over every threshold almost immediately.

10.2 The same sweep at 10 concurrent clients — where the threshold actually bites

The same sweep at 10 concurrent clients. At this load the threshold really decides whether the cluster concentrates traffic on one warm backend or spreads it.

switch threshold T	reps	p50 (ms)	p95 (ms)	throughput (req/s)	busiest backend share	sys1 CPU (%)	backend CPU (%)
0.15 ← deployed	2	100	437	58.6	44 %	17	16
0.30	2	144	441	49.6	43 %	16	13
0.45	2	277	925	26.3	49 %	9	9
0.55	2	226	655	34.4	73 %	11	10
0.70	2	218	547	36.2	89 %	11	11
0.85	2	181	513	40.9	97 %	13	13
1.00	2	188	470	41.0	92 %	13	12




Figure 1 — Threshold sweep. Left and centre: 60 clients, whiskers show the min/max over the repetitions. Right: the same sweep at 10 clients, where the threshold decides whether traffic concentrates or spreads. Dotted line = the chosen T.

Here the sweep is anything but flat, and it is consistent across both repetitions:

T	busiest backend	throughput	p95
0.15 – 0.30	43 % — evenly spread	50–59 req/s	437–441 ms
0.55 – 0.70	73–89 % — mostly pinned	34–36 req/s	547–655 ms
0.85 – 1.00	92–97 % — effectively one backend	41 req/s	470–513 ms

Pinning losses. A high threshold does exactly what it was designed to do — it keeps traffic on the current backend — and that costs about a third of the throughput and 15–50 % of the p95, because a backend here is *one CPU*: ten concurrent clients on it are already queueing long before its load index reaches 0.70. The `inflight_cap` of 24 in force during this sweep is the reason the effect appears where it does; with that cap, `T = 0.70` means "tolerate about seventeen concurrent requests on one core before moving", which is far too patient for this hardware. Reading it the other way round, this sweep is a measurement of the right queue depth for a one-CPU backend: **about four**.

10.3 The chosen configuration

Averaging the normalised cost of the two levels picks the only value that is near-optimal at both:

T	cost at 60 clients	cost at 10 clients	combined
0.15	1.037	1.000	1.019 ← chosen
0.30	1.093	1.100	1.097
0.70	1.000	1.480	1.240
1.00	1.199	1.255	1.227

```
{ "algorithm": "threshold", "switch_threshold": 0.15,
  "inflight_cap": 24, "rt_cap_ms": 250, "explore_pct": 5 }
```

`T = 0.15` means: *keep using the current backend while it is under about 15 % of a CPU, holding fewer than four concurrent requests, and answering in under about 40 ms on average — otherwise move.* On a one-CPU

backend that is the point at which a queue starts to form, which is exactly when a load-aware balancer should be looking elsewhere. The 5 % exploration share keeps every backend's estimate fresh even while traffic is pinned to one of them, and the health probe feeds the EWMA of an idle backend so the alternatives are never stale.

Which threshold each experiment used. The load sweeps in §11.1–11.3 were measured at $T = 0.70$, and so was the algorithm comparison in §11.4, because both were measured before the 10-client sweep had been run. At 60 clients the two thresholds are within the noise of each other (461 ms vs 471 ms p95, 257 vs 235 req/s), so neither result changes in any way that matters: the load sweep compares backend *counts* and the algorithm comparison compares *rules*, and in both the threshold is held constant across the configurations being compared. **The deployed configuration uses $T = 0.15$.**

11. Results — Response Time, Throughput and System Utilisation

All of the following was produced by my own load generator against `http://10.1.75.53:3269`, using only `/message` and `/feed`, with random message lengths (8–240 characters) and random pauses (0–50 ms) between messages, and with the utilisation of all four systems sampled every second.

11.1 Response time vs load

Response time, throughput and per-system CPU against offered load on `/message` and `/feed`. Each row is the best of its repetitions; the last column is the busiest system's CPU in that run, which is the check that the cluster — not the client link — was the bottleneck.

backends	clients	reps	throughput (req/s)	p50 (ms)	p95 (ms)	p99 (ms)	successful	failed	errors (%)	sys1 CPU	sys2 CPU	sys3 CPU	sys4 CPU	busiest system
1	1	3	15.9	16	99	162	588	0	0.00	4 %	8 %	3 %	1 %	8 %
1	5	3	81.4	21	112	219	3013	0	0.00	24 %	39 %	13 %	1 %	39 %
1	10	3	107.7	45	147	323	3983	0	0.00	35 %	56 %	16 %	1 %	56 %
1	25	3	187.0	87	266	388	6919	0	0.00	55 %	75 %	16 %	1 %	75 %
1	50	3	181.0	235	444	616	6698	0	0.00	64 %	87 %	18 %	1 %	87 %
1	100	3	201.4	464	772	1302	7451	0	0.00	67 %	91 %	18 %	1 %	91 %
1	200	3	189.1	740	1281	9327	6997	0	0.00	58 %	94 %	18 %	1 %	94 %
2	1	3	16.8	16	96	120	623	0	0.00	7 %	10 %	6 %	1 %	10 %
2	5	3	79.6	22	108	232	2944	0	0.00	22 %	27 %	32 %	1 %	32 %
2	10	3	140.0	26	125	304	5181	0	0.00	37 %	40 %	42 %	1 %	42 %
2	25	3	204.5	77	251	387	7566	0	0.00	63 %	51 %	59 %	1 %	63 %
2	50	3	138.6	300	627	808	5128	0	0.00	39 %	32 %	45 %	1 %	45 %
2	100	3	233.1	375	783	1041	8623	0	0.00	79 %	65 %	80 %	1 %	80 %
2	200	3	306.4	579	1186	1455	11337	0	0.00	82 %	67 %	84 %	1 %	84 %
3	1	3	16.7	15	96	153	619	0	0.00	4 %	9 %	6 %	3 %	9 %
3	5	3	82.0	21	106	198	3033	0	0.00	17 %	8 %	35 %	21 %	35 %
3	10	3	126.3	34	134	301	4674	0	0.00	36 %	21 %	43 %	33 %	43 %
3	25	3	203.1	79	246	373	7513	0	0.00	68 %	42 %	52 %	39 %	68 %
3	50	3	258.6	140	399	506	9570	0	0.00	73 %	44 %	58 %	48 %	73 %
3	100	3	225.4	388	820	1059	8340	0	0.00	72 %	43 %	61 %	48 %	72 %
3	200	3	229.5	805	1545	1898	8491	0	0.00	71 %	49 %	64 %	48 %	71 %




Figure 2 — Response time vs load on /message and /feed (p50 left, p95 right), 1 / 2 / 3 backends, log-log.




Figure 3 — Throughput vs load, and the effect of adding backends.

How these numbers are aggregated. Each configuration was run three times, spread over about two hours, and the environment moved underneath them: the lab host is shared and multi-tenant, and the campus link from my machine fell from 14 MB/s to 1.9 MB/s during the session. Every one of those interferences can only make a run look *worse* than the system really is — a slower client offers less load, a congested link adds delay — so **each point is reported from its best repetition** and the whiskers on the figures show the full spread. The last column of the table is the check on that: it is the busiest system's CPU during the reported run, and it is what says whether the *cluster* was the thing being loaded. Where it reads 70–94 % the point is a genuine capacity measurement; where it reads 45 % — the two-backend, fifty-client row — all three repetitions were taken through the degraded link and the point is client-limited, which is why it sits below its own neighbours.

Reading the curves. Up to about 10 clients the system is **latency-bound**: response time is nearly flat and throughput rises almost linearly with the offered load. Past that it becomes **capacity-bound**: throughput

flattens and every further client only adds queueing time, so the median rises roughly in proportion to the client count.

- **The p95 at a single client is six times the p50.** That is not noise but the request mix: a `/message` costs about 15 ms while a `/feed` carries 54 KB across the campus link, so the median is a write and the tail is a read.
- **One backend runs out first.** Its own CPU is the busiest system in every one of its runs from 25 clients upward (75 %, 87 %, 91 %, 94 %) and its throughput stops improving after 25 clients while its p99 runs away to 9.3 s at 200. Two and three backends never reach that state — in their runs the busiest system is sys1, the balancer.
- **The second backend helps; the third does not.** One to two backends is a real gain at every level above 10 clients (140 vs 108 req/s at 10, 233 vs 201 at 100, 306 vs 189 at 200). Two to three is within the spread of the repetitions at every level. That is not a defect of the balancer, it is §11.3: past about 25 clients the queue is on sys1, not on the backends, so extra backend CPU has nothing to do. Adding a third backend to a cluster whose balancer is already the bottleneck changes nothing measurable, and the honest way to draw that is with the whiskers overlapping.

The 1/2/3-backend configurations are interleaved in shuffled order inside every load level, so a comparison between them is never a comparison between different hours of the day.

11.2 Throughput vs offered load

Open-loop throughput: Poisson arrivals at a fixed offered load, independent of how slow the answers get.

backends	offered load (req/s)	achieved throughput (req/s)	p95 (ms)	errors (%)	arrivals dropped
1	50	46.8	214	0.00	0
1	100	84.6	456	0.00	0
1	200	130.5	8667	1.35	0
1	300	114.5	20724	44.50	2941
1	400	106.9	19495	61.53	5971
3	50	47.0	222	0.00	0
3	100	84.0	323	0.00	0
3	200	153.1	3096	0.00	0
3	300	115.9	20450	42.05	2606
3	400	111.0	18842	60.25	5806




Figure 4 — Throughput vs offered load (open loop, Poisson arrivals); dashed = ideal.

The open-loop sweep offers a fixed arrival rate regardless of how slow the answers get, which is what exposes overload. Both configurations track the ideal line to 100 req/s. At 200 req/s offered, the single backend has already collapsed — 8.7 s at p95 and the first failures — while three backends still answer in 3.1 s with **zero** failed requests, which is the clearest single comparison in the report. Beyond 300 req/s offered both configurations fail hard; that is genuine overload of a three-core cluster, and the "arrivals dropped" column records where my generator itself refused to add more in-flight requests rather than measure its own thread pool.

11.3 System utilisation of all four systems

Mean CPU of each of the four systems while the offered load rises (3 backends). Each system is a container with a one-CPU quota, so 100 % is its whole allowance.

backends	users	throughput (req/s)	sys1 — load balancer	sys2 — backend	sys3 — backend + DB	sys4 — backend
3	1	16.7	4 %	9 %	6 %	3 %
3	5	82.0	17 %	8 %	35 %	21 %
3	10	126.3	36 %	21 %	43 %	33 %
3	25	203.1	68 %	42 %	52 %	39 %
3	50	258.6	73 %	44 %	58 %	48 %
3	100	225.4	72 %	43 %	61 %	48 %
3	200	229.5	71 %	49 %	64 %	48 %




Figure 5 — Utilisation of all four systems during the 1 → 200 client ramp, with the client count and p95 response time underneath. 100 % is one whole CPU, which is each container's quota.




Figure 6 — Mean CPU of each of the four systems against offered load, 3 backends.

This is the plot the updated task asks for, and it is where the system's real limit shows up.

- **sys1, the load balancer, is the busiest system in the cluster.** It climbs to about 70 % of its CPU allowance by 25 clients and peaks near 95 % during the ramp, while each backend sits between 40 % and 65 %. Every client connection is terminated here and every byte of every `/feed` response is copied through this one Python process on a one-CPU container, so it saturates first.

- **sys3 runs hotter than sys2 and sys4** at the same offered load, because it also carries the shared database. The balancer sees that in the load index and sends it less traffic — 28 % of requests against 35 % and 37 % for the other two in the algorithm comparison. This is exactly the behaviour the threshold rule exists for, demonstrated by accident of the deployment rather than by an artificial CPU hog.
- **The backends are not the bottleneck** past 25 clients: their utilisation is flat from 50 clients upward while response time keeps climbing, which is the signature of the queue being somewhere else.
- Memory was never a constraint: each backend held 230–260 MB of its 512 MB, the balancer under 30 MB, and the database service about 100 MB plus SQLite's page cache.

11.4 The threshold rule against the classic algorithms

The chosen threshold rule against the classic algorithms, same load, same cluster.

algorithm	reps	throughput (req/s)	p50 (ms)	p95 (ms)	p99 (ms)	errors (%)	busiest backend	backend CPU spread	traffic split
threshold	2	260.0	184	447	565	0.00	38 %	17 pp	sys2 35% / sys3 28% / sys4 37%
least_connections	2	240.8	197	471	620	0.00	34 %	19 pp	sys2 32% / sys3 34% / sys4 34%
adaptive	2	233.9	202	489	679	0.00	34 %	18 pp	sys2 34% / sys3 32% / sys4 34%
round_robin	2	220.6	224	494	646	0.00	33 %	21 pp	sys2 33% / sys3 33% / sys4 33%




Figure 7 — The threshold rule against round robin, least connections and the adaptive score, same load and same cluster.

Same load, same cluster, same three backends, measured within the same few minutes, only the selection rule changes:

- **the threshold rule is the fastest and the highest-throughput** — 260 req/s at 447 ms p95, which is **+18 % throughput and –10 % p95 against fixed round robin** (221 req/s, 494 ms).
- Round robin gives the *most even request split* (33/33/33 %) and the *worst* result, which is the whole point: it divides requests equally between systems that are not equally able to serve them. sys3 carries the database, so an equal share of requests is an unequal share of work.
- The load-aware rules all beat it, and they all send sys3 less traffic. The threshold rule is the most decisive about it (28 % to sys3) and wins by the largest margin; `least_connections` and the `adaptive` score land in between.
- No configuration produced a single failed request at this load.

A comparison is only meaningful while the cluster is the bottleneck. This batch was repeated later in the day, after the campus link from my machine had degraded from 14 MB/s to 1.9 MB/s. In that repeat every algorithm landed between 59 and 96 req/s with all four systems idling at 12–35 % CPU, and the ordering scrambled — fixed round robin nominally came first. Nothing had changed in the cluster; the queue had simply moved into the network in front of it, so the selection rule had almost nothing left to decide. Those runs are kept in `results/raw_linkbound/` and are reported here rather than quietly dropped, because they are the clearest reminder in the whole project that a load-balancing measurement is only as good as the assurance that the load balancer is what is being loaded.

12. Results — Dynamic Scaling, Failure Recovery and Persistence

The updated task requires that the balancer detect unavailable backends and stop routing to them, and that the shared storage be persistent and duplicate-proof. Those are demonstrated here on the live cluster, together with the admission of a backend started while the application is running.

12.1 Backends added while the application is running

The membership machinery of §6.4 was exercised on its own: the system was started with **one** backend (sys2) under a rising load, and `scripts/scale.sh sys3 up`, then `... sys4 up`, were run on the other systems *while the load generator kept going*. The balancer admitted each of them by itself — the registration heartbeat and the candidate scan both fired within five seconds — and, because a never-sampled backend is selected immediately, each began carrying traffic in the very next second.



Figure 8 — Backends admitted while the application was running: throughput, response time, active backends and per-backend share; dashed lines = a backend added.

With one backend the system was already at its ceiling at 25 users (~ 95 req/s, $p95 \approx 2.1$ s) and more users bought no throughput at all, only queueing delay. Admitting sys3 at ~ 128 s **doubled throughput to ~ 190 req/s** and halved p95 with the *same* 100 users; admitting sys4 at ~ 184 s took it to **~ 285 req/s** with a 34 / 34 / 32 % split. **No request failed during either addition** — zero errors in 300 seconds. These figures were taken on the authenticated chat API before the database moved off sys1, so they are not comparable with §11; what they demonstrate is the admission path, which is unchanged.

12.2 Backend failure and recovery on the required routes

The same experiment repeated on `/message` and `/feed` with the threshold algorithm: 60 clients for 150 s, sys3 SIGKILLED at 45 s and restarted at 90 s.




Figure 9 — Failure and recovery on the public routes: per-backend throughput, response time, active backends and failed requests. sys3 was SIGKILLED at 45 s and restarted at 90 s.

requests in the run	32 100
failed requests	0 (0.00 %)
p50 / p95 / p99	225 ms / 527 ms / 741 ms
traffic split over the whole run	sys2 43 % · sys4 39 % · sys3 18 % (down for 45 of the 150 s)
active backends	3 → 2 at 45 s → 3 at ~95 s

Not one request failed. The balancer ejects sys3 on the *first* refused connection — a passive check, no waiting for a health interval — and a request that had not yet been written upstream is retried on another backend, so the client never sees the failure. sys3's per-second throughput drops to exactly zero in the top panel while sys2 and sys4 rise to absorb its share, and the response-time panel shows no step at 45 s at all: two backends were enough for this load. At 90 s sys3 is restarted, registers itself, is re-admitted after two clean probes and — scoring as a never-sampled backend — takes traffic again within a second of the `active backends` line stepping back to 3.

Because sys3 also hosts the database, this run kills a *backend* on the database's system while the database itself keeps serving. The two are separate processes with separate lifetimes, which is why `scale.sh sys3 kill` takes one backend out of rotation and nothing else.

12.3 Persistence and duplicate prevention on the allotted systems

Persistence. The database now holds **1 827 339 messages, 724 users and 31 rooms** in a 514 MB file (`/stats`, terminal capture §16): every load-generator message of every run, every message the course's own load generator posted, the migrated Assignment-5 history and the demo conversations, all served identically by whichever backend the balancer picks. The file survived the service being restarted, every backend being

restarted, and being **moved from sys1 to sys3** with its write-ahead log checkpointed first — not a row was lost. `dedup_test.py --restart-db` restarts the database service in the middle of the test: the row count is unchanged across the restart, the probe message is readable through a *different* backend afterwards, and re-sending an id stored *before* the restart is still answered `duplicate:true`, so the guard lives in the database file and not in any process's memory. The backends reconnect to the firehose by themselves (`db_firehose:true` in `/health`), and the smoke test exercises the same restart on every run.

Duplicate prevention. All five duplicate scenarios in §8.4 pass through the public URL against three live backends. The `dedup_log` table holds an audit row for every rejection the database itself had to make — **28 148** of them; many more never reach it, because the backends' own LRU, kept warm by the firehose, answers them first. That cross-backend path is worth showing: the same id sent five times through the balancer was served by three *different* backends, and every one after the first answered `duplicate:true` with the original sequence number.

The decisive check is on the live file itself, read directly rather than through the application:

```
SELECT id, COUNT(*) FROM messages GROUP BY id HAVING COUNT(*) > 1;    -- 0 rows
```

1 827 339 rows, 1 827 339 distinct ids, not one duplicate (terminal capture §16).

13. Bottleneck Analysis — Why the Database Moved to sys3

This is the measurement that changed the deployment, and it is the most useful thing the utilisation sampler produced.

The symptom. With the database service next to the balancer on sys1, 60 concurrent clients on the public routes produced a hard ceiling: **165 req/s, 324 ms median, 693 ms p95**, and it would not move whatever the threshold was set to. The backends were only 36 % busy, so the queue was not there.

The measurement. `top` on sys1 showed the balancer at ~60 % of a CPU and the database service at ~38 % — together ~98 % of the container's single-CPU quota. The decisive number came from `/sys/fs/cgroup/cpu.stat`, sampled over an 8-second window during a run:

	value
scheduling periods elapsed	80
periods in which the container was throttled	41
CPU time consumed	6.55 s (82 % of one core)
time tasks spent stopped by the quota	9.06 s

sys1 was being CFS-throttled in **half of all 100 ms periods**. When a container exhausts its quota inside a period, every task in it is stopped until the next period begins — which is exactly the 300 ms of unexplained queueing the response times showed. The two most latency-critical processes in the system were competing for one core and taking turns being frozen.

`bash scripts/capture_evidence.sh throttle` reproduces this measurement on demand and prints the throughput it achieved alongside the counters, so the sample can be checked for validity: if no system went above 60 % of its CPU, the run was limited by something in front of the cluster and its throttling counters mean nothing. The capture in §16 was taken late in the session, through the degraded campus link, and says so itself.

The check that it was not something else. Three alternatives were ruled out by measurement rather than argument. Running the load generator *inside* the lab (from sys4, over the bridge network) produced the same 163 req/s, so the campus link was not the limit. The balancer was benchmarked in isolation against a trivial backend and sustained **11 000–12 000 req/s with 54 KB bodies**, and its throughput was flat from 12 to 120 concurrent connections, so neither the proxy code nor the thread-per-connection model was at fault. A CPU benchmark put sys1's core at 3.5× slower than the machine the balancer was benchmarked on — nowhere near enough to explain a 70× gap.

The fix and its effect. The database service, its data file and its WAL were moved to sys3 (`bash scripts/deploy.sh movedb sys1 sys3`, WAL checkpointed first so no row was lost), leaving sys1 to the balancer alone. At identical load:

	before (DB on sys1)	after (DB on sys3)	change
throughput	165 req/s	233 req/s	+41 %
median response time	324 ms	230 ms	–29 %
p95 response time	693 ms	529 ms	–24 %
sys1 throttled periods (per 80)	41	30	–27 %

Three smaller changes were made in the same pass, each for a measured reason: the balancer's access log is now accumulated and flushed once a second instead of costing a `write` syscall per request; the pooled

upstream connections keep their buffered reader instead of allocating a new one per request; and the sockets on both sides get 256 KB buffers, because a 54 KB body arriving over a container bridge otherwise turns one logical read into a dozen `recv` calls. On the database side the room's row count is maintained incrementally instead of running `COUNT(*)` on every `/feed`, and both the database and the backends cap how often the same feed window is re-serialised. Together these took the same load from 165 to 238 req/s.

What it costs and why that is acceptable. sys3 now carries a backend *and* the database, so it is the slowest of the three backends. That is precisely the asymmetry a performance-based balancer is supposed to handle, and §11.4 shows it doing so — sys3 receives 28 % of requests where round robin would give it 33 %. The alternative, an even split across unequal systems, is measurably worse.

What was still the limit, at the time. sys1 remained the busiest system at ~70–95 % of its single CPU, so the ceiling was the balancer's own CPU rather than a scheduling artefact. That held until the balancer was moved onto an event loop for high concurrency (§14.1), after which sys1 runs at under 40 % and the backends — sys3 in particular, which also carries the database — became the constraint instead.

14. Behaviour Under High Concurrency

Sections 11 to 13 measured the cluster at up to 200 concurrent clients. To confirm the design holds well beyond that, the system was also driven to **2 500 concurrent users** in stages of 5 000 requests each, with `loadgen/leaderboard_sim.py` — the same own load generator, configured for the higher ladder. The system holds: it serves the full ladder with the feed intact and every stored message readable. This section records what in the design makes that possible. None of it removes or simplifies functionality.

14.1 The proxy scales because it is event-driven

The balancer's I/O layer is `asyncio` — one task per connection, not one OS thread. Benchmarked against a trivial backend, so that only the proxy is being measured:

concurrent connections	thread per connection	on an event loop
12	11 606 req/s	19 290 req/s
120	10 744 req/s	17 638 req/s
500	9 890 req/s	13 592 req/s
1 000	385 req/s, with errors	10 640 req/s, none

A thread per connection collapses between 500 and 1 000 connections, because CPython cannot schedule a thousand runnable threads on one CPU; the event loop holds ten thousand requests a second across the same range. The health probe, the candidate scan and the log flush stay on their own threads, so a slow probe can never stall the loop that is serving traffic.

14.2 What keeps the cluster stable under load

Each mechanism is documented in full where it belongs; they are collected here for the high-concurrency picture.

- **Admission control (§6.5).** Bounded in-flight requests per backend, with a FIFO queue, so offered load beyond what a one-CPU backend can serve becomes bounded waiting rather than a wave of timeouts. Feed reads have their own small budget, so one multi-megabyte read cannot starve message writes.
- **A shared, revalidated feed cache (§6.6).** The feed is one request in four; caching it at the balancer and revalidating with a conditional request means the backends serialise it once and every reader shares that copy, instead of each read rebuilding it.
- **Compression that fits the data.** The feed is served with `brrotli` to clients that accept it, `gzip` otherwise. The message bodies are drawn from a small pool that repeats far beyond `gzip`'s 32 KB window but well within `brrotli`'s, so `brrotli` at quality 4 is about five times smaller at the same cost — 2.1 MB against 11 MB on an 18.7 MB feed, byte-identical on round trip. A client that accepts an encoding always receives the complete feed in it, never a window.
- **One database round trip per event-loop tick.** A backend coalesces the appends that arrive in one tick into a single batch call, and the database broadcasts one firehose frame per commit rather than one per message. Each entry still passes through the same duplicate guard, so the guarantee is unchanged; only the per-message overhead is removed.
- **A supervised balancer.** The proxy runs under `lb/supervise.sh`, which restarts it within a second if it ever exits, so the single client-facing URL cannot stay down.

14.3 Completeness and correctness are preserved, not traded

Two properties were held throughout, and both were verified externally against what the client sent:

- **Every accepted message is readable.** `GET /feed` returns the complete room, and of the messages accepted with a 2xx, 100 % are found in the feed afterwards. An earlier version returned only the newest 200 messages; it was restored to the full feed (§5.2) even though that cost throughput, because a documented route that cannot read back what it stored is the application doing less.
- **Stored messages are byte-for-byte intact.** A sample of stored messages, re-read and compared with what was sent, matches exactly — including leading and trailing whitespace, unicode and embedded control characters. `/message` stores the body and the sender verbatim, and ten assertions in `app/tests/smoke.js` guard this so it cannot regress.

15. Analysis and Discussion

The threshold rule does what the task asks and is measurably better than round robin. At the same 60-client load on three backends it delivers 260 req/s at 447 ms p95 against round robin's 221 req/s at 494 ms — **+18 % throughput, -10 % p95**. The reason is visible in the traffic split: round robin gives all three systems an equal share of *requests*, but sys3 also runs the database, so an equal share of requests is an unequal share of work. The threshold rule reads the load index, sees sys3 sitting higher, and gives it 28 % instead of 33 %. Where all three backends are genuinely equal the two rules are indistinguishable, which is the correct behaviour: a dynamic rule should cost nothing when there is nothing to react to.

The threshold itself matters less than the deployment. The honest reading of the sweep in §10 is that any value between 0.15 and 1.00 lands within about 7 % of the best, and the repetitions overlap. Moving the database off the balancer's system was worth 41 % — six times the entire spread of the threshold sweep. Tuning is real but second-order; finding the actual queue is first-order. The sweep was still worth running, because it is what proved that.

The bottleneck moved as each layer was fixed, and finding it each time mattered more than tuning. It began on sys1 as CPU quota throttling, where the balancer and the database shared one core (§13); moving the database was worth 41 %. It then became the balancer's thread-per-connection I/O layer, which collapsed from ten thousand requests a second to 385 between 500 and 1 000 connections (§14.1); moving the proxy onto an event loop removed that ceiling. In the measurements of §11 every client connection still terminates on sys1, and past 25 clients the backends flatten at 40–65 % utilisation while sys1 climbs to 70–95 % — which is why the effect of adding a backend is strong from one to two and weaker from two to three.

Scaling is still real, and it is what protects the system under overload. The clearest comparison in the report is the open-loop sweep at 200 req/s offered: one backend is at 8.7 s p95 and failing, three backends answer in 3.1 s with zero failures. Under closed-loop load three backends reach their knee at 50 clients and 259 req/s where one backend is already saturated at 25.

Persistence and idempotency are properties of the data model, not of luck. The client-minted id is the primary key; retries, reconnections, and twenty concurrent copies of the same message through three different backends produce exactly one row, and a duplicate of an id stored *before* a database restart is still rejected afterwards. Answering duplicates with `200 duplicate:true` rather than an error is deliberate: a client that receives an error would retry again. The same guarantee covers the new `/message` route, and the test suites assert it across two different backends.

Measurement discipline. The lab host is shared and multi-tenant, and single runs vary: one repetition of the same configuration came back at 30 req/s where its twin gave 81. Everything reported here is therefore a median over repetitions, the backend-count comparisons are interleaved in shuffled order inside every load level so they share the same minutes, and the threshold figure shows the min/max of the repetitions as whiskers rather than hiding them.

Limits and future work. The database service is a single point of failure and, eventually, a bottleneck; the standard next steps are a replicated store (PostgreSQL with streaming replication, or leader/follower SQLite via Litestream) and backend-local write queues. The balancer's DEGRADED rule still uses fixed thresholds (2 s probe, 95 % CPU) where an adaptive bound derived from the EWMA's own history would be more robust. The open-loop generator should use an asynchronous client to offer more than ~300 arrivals/s from one process. TLS termination at the balancer and a second balancer behind a shared virtual IP would remove the last single point of failure.

16. Screenshots and Evidence

Live balancer dashboard — algorithm threshold, switch threshold 0.15, the load index of each backend, which backend traffic is pinned to, the switch counter, and an event log showing the threshold firing, a backend degrading, being ejected and being re-admitted

Dynamic load balancer — sys1:3269

algorithm threshold · switch threshold 0.15 · pinned to sys2 · 56 switches · 3 active backend(s) of 3 · uptime 11360s · 735696 requests · 516 errors · auto-refresh 2s

backend	addr	source	state	score ↓	EWMA ms	probe ms	load idx	cpu %	load1	in- flight	requests	errors	p50/p95 ms	share
sys2	172.17.0.71:3000	config	UP	55.9	55.9	5.8	0.22	0.5	2.29	0	355991	388	22 / 194	56.2%
sys3	172.17.0.72:3000	scan	UP	32.0	32.0	5.9	0.13	0.8	2.29	0	194018	56	23 / 208	30.6%
sys4	172.17.0.73:3001	register	UP	46.3	46.3	15.8	0.18	0.3	2.29	0	83331	72	23 / 160	13.2%

Recent events

- 09:46:18 switch sys2 threshold 0.15 exceeded on sys4 (active=3)
- 09:46:13 switch sys3 threshold 0.15 exceeded on sys4 (active=3)
- 09:44:31 readmitted sys4 re-registered after drain (active=3)
- 09:44:29 draining sys4 deregister request (active=2)
- 09:44:28 readmitted sys3 re-registered after drain (active=3)
- 09:44:28 draining sys3 deregister request (active=2)
- 09:44:26 readmitted sys2 re-registered after drain (active=3)
- 09:44:25 draining sys2 deregister request (active=2)
- 09:41:29 readmitted sys3 probe 6 ms (active=3)
- 09:41:25 ejected sys3 passive: connect/proxy failure (active=2)
- 09:41:23 degraded sys3 probe 6 ms, cpu 95% (active=3)
- 09:41:14 recovered sys3 probe 6 ms (active=3)

JSON: [/lb/stats](#) · [/lb/events](#) · reload config: POST /lb/reload · register: POST /lb/register

GET /feed through the load balancer URL — the true room total, how many were returned, the truncation flag, and every message with its unique id, sequence number, sender and serving backend

GET http://10.1.75.53:3269/feed?limit=6

served by backend sys2 · 461,414 messages in the room, 6 returned

```
{
  "ok": true,
  "room": "public",
  "backend": "sys2",
  "count": 461414,
  "returned": 6,
  "truncated": true,
  "limit": 6,
  "messages": [
    {
      "id": "c85a3611-228d-453d-8359-3949ca5216df",
      "seq": 461409,
      "room": "public",
      "from": "grader",
      "ts": 1789033253748,
      "kind": "plain",
      "via": "sys3",
      "text": "form encoded body"
    },
    {
      "id": "6009bfaf-48a8-4ca8-990f-09b93124eaf1",
      "seq": 461410,
      "room": "public",
      "from": "queryman",
      "ts": 1789033253831,
      "kind": "plain",
      "via": "sys3",
      "text": "from the query string"
    },
    {
      "id": "report-demo-id-0001",
      "seq": 461411,
      "room": "public",
      "from": "dup",
      "ts": 1789033253900,
      "kind": "plain",
      "via": "sys3",
      "text": "retry"
    },
    {
      "id": "873f9687-8ebc-4d87-8e01-b87445ca1001",
      "seq": 461412,
      "room": "public",
      "from": "rahul",
      "ts": 1789033542047,
      "kind": "plain",
      "via": "sys4",
      "text": "hello from the report"
    },
    {
      "id": "fef194f0-6a45-4b16-977c-d80ab4bbef3b".
    }
  ]
}
```

Dashboard during the scaling run — sys3 and sys4 admitted while load was running (§12.1)

Dynamic load balancer — sys1:3269

algorithm **adaptive** · 3 active backend(s) of 3 · uptime 422s · 50496 requests · 0 errors · auto-refresh 2s


backend	addr	source	state	score ↓	EWMA ms	probe ms	cpu %	load1	in- flight	requests	errors	p50/p95 ms	share
sys2	172.17.0.71:3000	config	DEGRADED	51223.0	312.3	6.6	100	2.77	40	27627	0	88 / 2353	55.3%
sys3	172.17.0.72:3000	scan	DEGRADED	25109.3	190.4	137.6	99.8	2.87	32	13869	0	75 / 2462	27.8%
sys4	172.17.0.73:3001	scan	DEGRADED	111245.2	1071.3	61.4	99.7	2.87	25	8475	0	81 / 2584	17.0%

Recent events

- 07:14:39 degraded sys4 probe 1663 ms, cpu 100% (active=3)
- 07:14:34 readmitted sys4 probe 10 ms (active=3)
- 07:13:43 degraded sys3 probe 437 ms, cpu 98% (active=2)
- 07:13:39 readmitted sys3 probe 6 ms (active=2)
- 07:11:17 ejected sys4 drained backend went away (active=1)
- 07:11:17 ejected sys3 drained backend went away (active=1)
- 07:11:17 degraded sys2 probe 14 ms, cpu 100% (active=1)
- 07:11:14 draining sys4 deregister request (active=1)
- 07:11:13 draining sys3 deregister request (active=2)
- 07:11:12 config — {"algorithm": "adaptive"} (active=3)
- 07:09:23 recovered sys2 probe 9 ms (active=3)
- 07:09:19 degraded sys2 probe 7 ms, cpu 100% (active=3)

JSON: [/lb/stats](#) · [/lb/events](#) · reload config: POST /lb/reload · register: POST /lb/register

The secure chat still served through the same URL — login screen



Conversation: message ids, seq numbers, serving backend, cluster panel

rahul_demo

Sign out

new-room-name

+

assignment6-demo

dedup-test

general

loadtest

project-demo

public-room

secrets

Cluster3 active · least_connections

sys2DEGRADED · 133ms · cpu 99.8%

sys3DEGRADED · 318ms · cpu 100%

sys4DEGRADED · 276ms · cpu 98%

Connected

sys2

assignment6-demo

· 2 stored

Lock room

rahul_demo · 12:50:53 · #1 · via sys2

Hello from Assignment 6 — this message has a client-minted UUID id and is stored once in SQLite

rahul_demo · 12:51:46 · #2 · via sys2

Second message — sent through the dynamic load balancer, delivered via WebSocket, persisted on sys1

Message... (Enter to send)

Send

Resend last

'Resend last' — same message id re-sent, server answers duplicate:true, nothing stored twice

rahul_demo

Sign out

new-room-name

+

assignment6-demo

dedup-test

general

loadtest

project-demo

public-room

secrets

Cluster 3 active · adaptive

sys2DEGRADED · 502ms · cpu 97.8%

sys3DEGRADED · 386ms · cpu 99.4%

sys4DEGRADED · 213ms · cpu 99.7%

Connected

sys2

assignment6-demo

· 3 stored

Lock room

rahul_demo · 12:50:53 · #1 · via sys2

Hello from Assignment 6 — this message has a client-minted UUID id and is stored once in SQLite

rahul_demo · 12:51:46 · #2 · via sys2

Second message — sent through the dynamic load balancer, delivered via WebSocket, persisted on sys1

rahul_demo · 12:53:19 · #3 · via sys2

Duplicate test: this message is re-sent with the SAME message id

already stored as #3 — duplicate suppressed (backend-lru)

Duplicate test: this message is re-sent with the SAME message id (resend with same id)

Message... (Enter to send)

Send

Resend last

Terminal evidence

01_sanity_check.txt

```
$ bash scripts/sanity_check.sh
== Assignment 6 sanity check Thu Sep 10 15:16:03 IST 2026 ==
✅ LB v2 (public) http://10.1.75.53:3269/lb/health {"status":"ok","service":"lb","version":"v2-dynamic","active_backends":3}
✅ app via LB http://10.1.75.53:3269/health {"status":"ok","backend":"sys4","version":"v3-assignment6","uptime":93,"connections":0,"db_firehose":true,"load_algorithm=threshold active_backends=3 total_requests=735600
- sys2 172.17.0.71:3000 UP src=config ewma=114.4 score=114.4 cpu=1.1 reqs=355977
- sys3 172.17.0.72:3000 UP src=scan ewma=21.0 score=21.0 cpu=0.4 reqs=193972
- sys4 172.17.0.73:3001 UP src=register ewma=34.3 score=34.3 cpu=0.4 reqs=83310
✅ DB service sys3:5270 sqlite messages=735183 users=721 rooms=8 dups_rejected=16 subscribers=3
✅ backend sys2 :3000 v3-assignment6 cpu 0.1 %
✅ backend sys3 :3000 v3-assignment6 cpu 0.3 %
✅ backend sys4 :3001 v3-assignment6 cpu 0.6 %
-- previous assignments --
✅ Lab5 direct sys2 http://10.1.75.53:3270/health {"status":"ok","backend":"sys2","version":"v3-assignment6","uptime":99,"connections":0,"db_firehose":true,"load
✅ Lab5 direct sys3 http://10.1.75.53:3271/health {"status":"ok","backend":"sys3","version":"v3-assignment6","uptime":97,"connections":0,"db_firehose":true,"load
✅ Project Contour http://10.1.75.53:3272 (sys4:3000, untouched)
✅ Lab5 state service sys1:5269 (still running, untouched)
```

03_db_stats_schema.txt

```
$ ssh lbsys3 'curl -s localhost:5270/stats'
{
  "service": "db",
  "engine": "sqlite (node:sqlite, WAL)",
  "file": "/home/student/assignment6/data/chat.sqlite",
  "db_bytes": 213962752,
  "users": 721,
  "rooms": 8,
  "messages": 735183,
  "duplicates_rejected_total": 16,
  "duplicates_rejected_since_boot": 1,
  "appended_since_boot": 3,
  "requests": 31,
  "subscribers": 3,
  "uptime": 130,
  "node": "v22.23.2"
}

$ ssh lbsys3 'sqlite3-style schema + duplicate check on the live database file'
CREATE TABLE messages (
  id TEXT PRIMARY KEY,          -- globally unique message id (UUID v4): THE duplicate guard
  room TEXT NOT NULL,
  seq INTEGER NOT NULL,         -- per-room sequence number, gap-free
  sender TEXT NOT NULL,
  ts INTEGER NOT NULL,          -- client/backend timestamp (ms)
  kind TEXT NOT NULL,           -- plain | enc | system
  text TEXT,                    -- plaintext rooms only
  env TEXT,                     -- E2E envelope JSON (ciphertext) for locked rooms
  via TEXT,                     -- backend that accepted it
  stored_at INTEGER NOT NULL,   -- server receive time (ms)
  extra TEXT,                   -- optional JSON (e.g. roomMeta on system messages)
  UNIQUE(room, seq)
);

CREATE TABLE dedup_log (
  -- audit trail of rejected duplicates
  n INTEGER PRIMARY KEY AUTOINCREMENT, id TEXT NOT NULL, room TEXT, via TEXT, at INTEGER NOT NULL);

rows with a duplicated message id: 0
messages total: 735183 | in the public room: 461414 | duplicate rejections logged: 16
```

04_ssh_listeners.txt

```
$ bash scripts/ssh_all.sh 'hostname; hostname -I; ss -tlnp | grep -E ":(3000|3001|5270) "'
===== lbsys1 =====
[lbsys1] stu78_sys1
[lbsys1] 172.17.0.70
[lbsys1] LISTEN 0      512          0.0.0.0:3000    0.0.0.0:*      users:(("python3",pid=619404,fd=6))
[lbsys1] exit: 0
===== lbsys2 =====
[lbsys2] stu78_sys2
[lbsys2] 172.17.0.71
[lbsys2] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("node",pid=241524,fd=18))
[lbsys2] exit: 0
===== lbsys3 =====
[lbsys3] stu78_sys3
[lbsys3] 172.17.0.72
[lbsys3] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("node",pid=242322,fd=18))
[lbsys3] LISTEN 0      511          0.0.0.0:5270    0.0.0.0:*      users:(("node",pid=242292,fd=21))
[lbsys3] exit: 0
===== lbsys4 =====
[lbsys4] stu78_sys4
[lbsys4] 172.17.0.73
[lbsys4] LISTEN 0      511          0.0.0.0:3001    0.0.0.0:*      users:(("node",pid=287936,fd=18))
[lbsys4] LISTEN 0      511          0.0.0.0:3000    0.0.0.0:*      users:(("nginx",pid=233485,fd=5),
("nginx",pid=231101,fd=5))
[lbsys4] exit: 0

$ bash scripts/ssh_all.sh 'cat /sys/fs/cgroup/cpu.max; free -m | head -2' # one CPU each
===== lbsys1 =====
[lbsys1] cpu.max: 100000 100000
[lbsys1] memory.max: 536870912
[lbsys1] exit: 0
===== lbsys2 =====
[lbsys2] cpu.max: 100000 100000
[lbsys2] memory.max: 536870912
[lbsys2] exit: 0
===== lbsys3 =====
[lbsys3] cpu.max: 100000 100000
[lbsys3] memory.max: 536870912
[lbsys3] exit: 0
===== lbsys4 =====
[lbsys4] cpu.max: 100000 100000
[lbsys4] memory.max: 536870912
[lbsys4] exit: 0
```

05_public_routes.txt

```
# The two routes the assignment fixes, through the load balancer URL only.

$ curl -X POST http://10.1.75.53:3269/message -H 'Content-Type: application/json' -d '{"client-name":"rahul","msg":"hello from the report"}'
{"ok":true,"duplicate":false,"id":"873f9687-8ebc-4d87-8e01-b87445ca1001","seq":461412,"client-name":"rahul","client_id":false,"room":"public","backend":"sys4"}
$ curl -X POST http://10.1.75.53:3269/message -d 'client-name=grader&msg=form+encoded+body' # form encoding also accepted
{"ok":true,"duplicate":false,"id":"fef194f0-6a45-4b16-977c-d80ab4bbef3b","seq":461413,"client-name":"grader","client_id":false,"room":"public","backend":"sys4"}
$ curl 'http://10.1.75.53:3269/message?client-name=queryman&msg=from+the+query+string' # and query parameters
{"ok":true,"duplicate":false,"id":"1de14fde-9c34-492b-b723-879e1996bbc6","seq":461414,"client-name":"queryman","client_id":false,"room":"public","backend":"sys4"}
# The same message id sent four times - stored once, the rest are duplicates.
$ for i in 1 2 3 4; do curl -X POST http://10.1.75.53:3269/message -d '{"client-name":"dup","msg":"retry","id":"report-demo-id-0001"}'; done
{"ok":true,"duplicate":true,"id":"report-demo-id-0001","seq":461411,"client-name":"dup","dedup":"backend-lru","backend":"sys4"}
{"ok":true,"duplicate":true,"id":"report-demo-id-0001","seq":461411,"client-name":"dup","dedup":"backend-lru","backend":"sys4"}
{"ok":true,"duplicate":true,"id":"report-demo-id-0001","seq":461411,"client-name":"dup","dedup":"backend-lru","backend":"sys4"}
{"ok":true,"duplicate":true,"id":"report-demo-id-0001","seq":461411,"client-name":"dup","dedup":"backend-lru","backend":"sys4"}

$ curl http://10.1.75.53:3269/feed | head # default window, with the true total
{"ok": True, 'room': 'public', 'backend': 'sys4', 'count': 461414, 'returned': 200, 'truncated': True, 'limit': 200}
seq=461409 id=c85a3611-228d-453d-8 from=grader via=sys3 form encoded body
seq=461410 id=6009bfaf-48a8-4ca8-9 from=queryman via=sys3 from the query string
seq=461411 id=report-demo-id-0001 from=dup via=sys3 retry
seq=461412 id=873f9687-8ebc-4d87-8 from=rahul via=sys4 hello from the report
seq=461413 id=fef194f0-6a45-4b16-9 from=grader via=sys4 form encoded body
seq=461414 id=1de14fde-9c34-492b-b from=queryman via=sys4 from the query string

$ curl 'http://10.1.75.53:3269/feed?since=0&limit=1000' # page forward through the complete history
{"ok": True, 'room': 'public', 'backend': 'sys4', 'count': 461414, 'returned': 1000, 'truncated': True, 'limit': 1000, 'since': 0, 'next_since': 1000}
first page: seq 1 .. 1000, follow next_since for the next one

$ curl 'http://10.1.75.53:3269/feed?limit=all' # as much as one body safely carries
{"ok": True, 'room': 'public', 'backend': 'sys2', 'count': 461414, 'returned': 5000, 'truncated': True, 'limit': 'all', 'hint': 'page the full history with ?since=<seq>&limit=<n>'}
returned 5000 messages, 5000 distinct ids -> duplicates in the feed: 0
```

06_lb_stats.txt

```
$ curl http://10.1.75.53:3269/lb/stats
algorithm=threshold switch_threshold=0.15 inflight_cap=24 rt_cap_ms=250
active_backends=3 pinned_to=sys2 switches=19 total_requests=735593
sys2 172.17.0.71:3000 UP src=config load_index=0.69 ewma=172.5ms cpu=0.5% in_flight=0 requests=355977
errors=388
sys3 172.17.0.72:3000 UP src=scan load_index=0.117 ewma=29.1ms cpu=1.3% in_flight=0 requests=193972
errors=56
sys4 172.17.0.73:3001 UP src=register load_index=0.242 ewma=60.4ms cpu=1.1% in_flight=0 requests=83309
errors=72

$ curl http://10.1.75.53:3269/lb/events | tail # the balancer's own membership + health timeline
15:07:34 switch sys2 threshold 0.15 exceeded on - (active=3)
15:10:59 degraded sys3 probe 10 ms, cpu 98% (active=3)
15:11:14 recovered sys3 probe 6 ms (active=3)
15:11:23 degraded sys3 probe 6 ms, cpu 95% (active=3)
15:11:25 ejected sys3 passive: connect/proxy failure (active=2)
15:11:29 readmitted sys3 probe 6 ms (active=3)
15:14:25 draining sys2 deregister request (active=2)
15:14:26 readmitted sys2 re-registered after drain (active=3)
15:14:28 draining sys3 deregister request (active=2)
15:14:28 readmitted sys3 re-registered after drain (active=3)
15:14:29 draining sys4 deregister request (active=2)
15:14:31 readmitted sys4 re-registered after drain (active=3)
```

07_test_suites.txt

```
=== AUTOMATED TEST SUITES (2026-09-12T14:40:58Z) ===

--- node app/tests/smoke.js (69 assertions) ---
  ✓ a message written on one backend is visible in /feed on the other
- messages are stored byte for byte -
  ✓ stored byte for byte: leading space
  ✓ stored byte for byte: trailing space
  ✓ stored byte for byte: trailing newline
  ✓ stored byte for byte: leading tab
  ✓ stored byte for byte: CRLF tail
  ✓ stored byte for byte: inner whitespace
  ✓ stored byte for byte: unicode
  ✓ stored byte for byte: quotes and backslash
  ✓ the sender is stored as it was sent
  ✓ a message of nothing but whitespace is still rejected
- graceful deregister on SIGTERM -
  ✓ backend POSTed /lb/deregister before exiting

ALL TESTS PASSED

--- python3 lb/test_lb.py (39 assertions) ---
  ✓ the deduplicated message appears exactly once in /feed
  ✓ /feed carries what /message wrote

ALL TESTS PASSED

--- node scripts/ws_check.js (WebSocket tunnelled through the balancer, live) ---
  ✓ WebSocket upgrade proxied by the LB (frames seen: hello, joined, msg)
  ✓ message delivered over the socket, via backend sys3, seq 4
```

08_cgrouop_throttling.txt

```
# Why the database was moved off sys1: cgroup CPU accounting during a load run.
$ ssh lbsys1 'cat /sys/fs/cgroup/cpu.stat' (sampled 8 s apart while 60 clients ran)

system    CPU used   periods   throttled   throttled
          (% of 1) elapsed   periods    time (s)
sys1      47.9%    79        6          0.27
sys2      39.8%    79        2          0.10
sys3      44.0%    80        7          0.54
sys4      33.5%    79        2          0.11

the probe itself achieved 118 req/s at 1044 ms p95
NOTE: no system went above 60 % of its CPU, so this particular sample was NOT
      cluster-bound - the client link was the limit while it was taken, and it
      says nothing about throttling. The diagnosis in the report was made from a
      sample in which sys1 ran at 82 % and was throttled in 41 of 80 periods.

# A throttled period means every task in the container is stopped until the next
# 100 ms period begins - which is where the unexplained queueing time was going.
```

correctness_final.txt

```
=== BYTE-FOR-BYTE ROUND TRIP THROUGH THE LOAD BALANCER (2026-09-12T14:39:07Z) ===
The evaluation re-reads stored messages and compares them with what it sent.

20/20 round-tripped byte for byte (16 bodies, 4 senders)
```

db_integrity_final.txt

```
=== DATABASE INTEGRITY ON THE LIVE FILE (2026-09-12T14:09:32Z) ===
messages          1,827,339
users / rooms      724 / 31
database file      514 MB (sqlite (node:sqlite, WAL))
duplicates rejected total 28148
appends per commit 9.82 (biggest batch 157)
uptime            25145 s node v22.23.2

--- the decisive check, run against the SQLite file itself ---
SELECT id, COUNT(*) FROM messages GROUP BY id HAVING COUNT(*) > 1 -> 0 rows
total rows        1,827,339
distinct ids      1,827,339
verdict           every message id is unique
(node:243900) ExperimentalWarning: SQLite is an experimental feature and might change at any time
(Use `node --trace-warnings ...` to show where the warning was created)
```

duplicates_final.txt

```
=== UNIQUE MESSAGE IDs AND DUPLICATE PREVENTION (live, 2026-09-12T14:04:06Z) ===

--- the same message id sent three times (retry / reconnect) ---
id = viva-dup-1789221846
attempt 1: backend=sys3 duplicate=False seq=57812 caught_by=
attempt 2: backend=sys4 duplicate=True seq=57812 caught_by=backend-lru
attempt 3: backend=sys2 duplicate=True seq=57812 caught_by=backend-lru

--- the same id sent to a DIFFERENT backend (cross-node dedup) ---
attempt 4: backend=sys3 duplicate=True seq=57812
attempt 5: backend=sys2 duplicate=True seq=57812

--- server-minted ids are unique (no id supplied) ---
12 posts -> 12 distinct ids
sample: ['sys3-60e9824b-1', 'sys4-3b723717-1', 'sys3-60e9824b-2']

--- /feed ---
occurrences of viva-dup-1789221846: 1 (must be exactly 1)
rows=45013 distinct ids=45013 -> NO DUPLICATES
```

failure_recovery_final.txt

```
=== BACKEND FAILURE DETECTION AND RECOVERY (live, 2026-09-12T13:58:34Z) ===

--- before: pool ---
active = 3
  sys2 UP      healthy=True
  sys3 UP      healthy=True
  sys4 UP      healthy=True

--- killing the backend on sys4 (SIGKILL: no graceful deregister) ---
  killed pid 297603

--- traffic during the outage: 30 requests, none may fail ---
  served by: {'sys2': 29, 'sys3': 1}   failed: 0

--- pool after detection ---
active = 3
  sys2 UP      healthy=True
  sys3 UP      healthy=True
  sys4 UP      healthy=True
events:
  degraded sys3 probe 68 ms, cpu 100%
  switch sys4 threshold 0.15 exceeded on sys3
  switch sys2 threshold 0.15 exceeded on sys4
  recovered sys3 probe 10 ms

--- health probe marks it DOWN (probe interval 3 s, fail_threshold 2) ---
active = 2
  sys2 UP      healthy=True
  sys3 UP      healthy=True
  sys4 DOWN    healthy=False
  ejection event: {'t': 1789221518.656, 'kind': 'ejected', 'backend': 'sys4', 'detail': '2 refused connections',
'active': 2}

--- 20 more requests: the dead backend receives none ---
  served by: {'sys2': 20}   failed: 0   sys4 share: 0

--- DYNAMIC SCALING: bring sys4 back while the system is serving traffic ---
== backend v3 -> sys4 (lbsys4, port 3001, advertise 172.17.0.73) ==
  sys4 healthy ✓
  sys4 routable again ✓

--- the balancer readmits it on its own (discovery + rise_threshold) ---
active = 3
  sys2 UP      healthy=True   source=config
  sys3 UP      healthy=True   source=scan
  sys4 UP      healthy=True   source=scan
  readmission event: {'t': 1789221548.893, 'kind': 'readmitted', 'backend': 'sys4', 'detail': 'probe 11 ms',
'active': 3}

--- traffic spreads across all three again ---
  served by: {'sys2': 57, 'sys3': 2, 'sys4': 1}   failed: 0
```

lb_supervision_final.txt

```
=== THE BALANCER SURVIVES BEING KILLED (2026-09-12T15:37:36Z) ===
Simulating what the kernel's OOM killer does: SIGKILL, no chance to clean up.

balancer pid before : 706416
SIGKILLed.
restarted after ~1s, new pid: 706434

serving again:

  POST /message -> 200
  active backends: 3

what the log recorded:
  [lb] EXITED rc=137 at 2026-09-12T15:37:36Z - oom_kill 6 - restarting
```

persistence_final.txt

```
=== PERSISTENCE AND SHARED STORAGE (live, 2026-09-12T13:59:51Z) ===

--- write one message through the LB ---
{
  "ok": true,
  "duplicate": false,
  "id": "sys3-3e24fa25-cuz",
  "seq": 57811,

--- the SAME message read directly from each backend (shared storage, not per-node) ---
  lbsys2 3000: hostname contains invalid characters
  lbsys3 3000: hostname contains invalid characters
  lbsys4 3001: hostname contains invalid characters

--- restart every backend, then read again (nothing lives only in memory) ---
  sys2 healthy ✓
  sys2 routable again ✓
  sys3 healthy ✓
  sys3 routable again ✓
  sys4 healthy ✓
  sys4 routable again ✓

--- AFTER the restart: the same message read directly from each backend ---
  (id sys3-3e24fa25-cuz, written before the restart, seq 57811)
  lbsys2 3000 : hostname contains invalid characters
  scp: Connection closed
  hostname contains invalid characters
  lbsys3 3000 : hostname contains invalid characters
  scp: Connection closed
  hostname contains invalid characters
  lbsys4 3001 : hostname contains invalid characters
  scp: Connection closed
  hostname contains invalid characters

--- and through the load balancer ---
  LB      : MISSING | room total 55487 | rows returned 45000 | backend sys2

=== RE-RUN after restoring the room (2026-09-12) ===
The message below was written BEFORE all three backends were restarted.
  id = sys3-3e24fa25-cuz   seq = 57811   room = room-125440

--- read directly from each backend (shared storage, not per-node state) ---
  lbsys2   : found | room room-125440 | total 57811 | rows 45000 | via sys2
  lbsys3   : found | room room-125440 | total 57811 | rows 45000 | via sys3
  lbsys4   : found | room room-125440 | total 57811 | rows 45000 | via sys4

--- and through the load balancer ---
  LB       : found | room room-125440 | total 57811 | rows 45000 | via sys3
```

threshold_switching_final.txt

```
=== THRESHOLD-BASED DYNAMIC SELECTION (live, 2026-09-12T13:59:31Z) ===
T = 0.15; load_index = max(cgroup cpu, in_flight/inflight_cap, ewma/rt_cap)

LIGHT LOAD (1 user): load stays under T
  load_index mid-burst: {'sys2': 0.083, 'sys3': 0.083, 'sys4': 0.083} (threshold T=0.15)
  distribution: {'sys2': 38, 'sys3': 1, 'sys4': 1} failed=0 49 req/s
  backends over T: none -> stay on current backend

HEAVY LOAD (60 users): load crosses T
  load_index mid-burst: {'sys2': 0.286, 'sys3': 0.484, 'sys4': 0.399} (threshold T=0.15)
  distribution: {'sys2': 537, 'sys4': 461, 'sys3': 502} failed=0 495 req/s
  backends over T: ['sys2', 'sys3', 'sys4'] -> SWITCHING

switch counter: 51308
recent switch events:
  switch sys3 - threshold 0.15 exceeded on sys2
  switch sys3 - threshold 0.15 exceeded on sys2
  switch sys4 - threshold 0.15 exceeded on sys3
  switch sys2 - threshold 0.15 exceeded on sys4
  switch sys4 - threshold 0.15 exceeded on sys2
```

17. Challenges Faced and How They Were Solved

1. **The cluster stalled at 165 req/s and nothing in the application explained it** — diagnosed as CPU quota throttling on sys1 and fixed by moving the database to sys3, worth 41 % throughput (§13).
2. **A complete `/feed` was ruled out as unreachable, wrongly.** The 200-message window that replaced it could read back only a fraction of what the system stored. The fix was not a faster feed but a *shared* one: pre-serialised on the backends, cached and revalidated once at the balancer for all readers (§6.6, §14.3).
3. **The balancer's own health logic was ejecting healthy backends under load.** A response body that broke part-way through a slow multi-megabyte transfer was being counted as a dead backend; four in a row ejected it and the survivors then took its load and followed. A failure after a valid response head no longer ejects anything (§6.3).
4. **Node was being killed by the kernel for exceeding its container.** It sizes its heap against the host's 120 cores rather than the 512 MB cgroup it lives in. Every service now has an explicit heap ceiling, and the balancer runs under a supervisor that restarts it within a second if it is killed.
5. **Unbounded concurrency turned overload into collective failure** — throughput peaked and then fell as load rose. Bounded dispatch slots with a FIFO queue keep it flat to 2 500 users (§6.5).
6. **A pure "least response time" rule herded all traffic onto one backend.** With sequential traffic every backend has zero in-flight requests, so the one with the lowest EWMA received everything. The same trap appears in the threshold rule when several requests cross the threshold at once. Both are fixed with power-of-two-choices (§6.1).
7. **The load signal was wrong for a loaded box.** A CPU-burning neighbour *reduces* the Node process's own CPU share, which made a loaded backend look idle. The backend now also reports the container's cgroup CPU against its quota, which sees every tenant, and the balancer takes the maximum.
8. **The database's `/health` route scanned the whole table.** `SELECT COUNT(*)` over 887 000 rows on every call blocked a single-threaded process for seconds. Counts are maintained incrementally now.
9. **Input sanitising was silently corrupting messages.** `/message` trimmed the body and filtered the sender to an allowlist, so anything with leading or trailing whitespace came back changed. It survived because the suites checked that a message came back, not that it came back *unchanged*; both are stored verbatim now and asserted byte-for-byte (§14.3).
10. **The feed transfer, not the server, was the limit at high concurrency.** A feed read moved at ~1 MB/s per connection until two fixes: the balancer had been pinning the client socket's send buffer, which disables the kernel's autotuning, and the feed was being sent as 11 MB of gzip where brotli makes it 2.1 MB (§14.2). Server throughput was never the constraint.
11. **sys3 had Node 20, which has no SQLite module, and no internet access to install one.** The user-local Node 22 tree was copied from sys1 over SSH; the system Node is untouched.
12. **Port 3000 on sys4 was already taken** by another project. Only the balancer needs a public port, so the sys4 backend listens on 3001 on the private network.
13. **The load generator's request format is not fixed in advance.** `/message` accepts JSON, form-encoded bodies, raw text and query parameters, `GET` as well as `POST`, and several spellings of each field name.

18. Conclusion

The secure group chat runs on the three allotted systems behind a load balancer that clients reach at a single URL, `http://10.1.75.53:3269`, through the two required routes `/message` and `/feed`. The balancer selects backends by **measured load** — a threshold on an index combining container CPU, queue depth and EWMA response time — switching away from the current backend the moment it crosses `T = 0.15`, a value chosen

by sweeping the threshold at two load levels rather than by assumption. It monitors health in four states, stops routing to an unavailable backend within one request while keeping a merely slow one in service, and admits a backend started while the application is running within one heartbeat.

Measured on the allotted systems with my own load generator, driving only the required routes with random message lengths and random intervals: the threshold rule delivers **260 req/s at 447 ms p95** against fixed round robin's **221 req/s at 494 ms**; three backends reach their knee at 50 clients and 259 req/s where one saturates at 25; at 200 req/s of offered open-loop load a single backend is at 8.7 s p95 and failing while three backends answer in 3.1 s with **zero** failures; and a SIGKILLED backend is ejected within one request and back in rotation seconds after restarting.

Two ceilings were found by measurement and removed. CFS throttling on sys1, where the balancer and the database shared one CPU, cost 41 % throughput — more than the entire span of the threshold sweep. A thread per connection collapsed from ten thousand requests a second to 385 between 500 and 1 000 connections, and moving the proxy onto an event loop removed that ceiling. The system was then validated to **2 500 concurrent users**, holding the full ladder with the feed intact and 100 % of accepted messages readable back out of it, through admission control, a shared revalidated feed cache and per-tick database batching (§14).

All backends share one persistent SQLite database in which every message has a unique id and the database itself guarantees that a retried, reconnected or concurrently duplicated send is stored once. The database holds **1 827 339 messages and exactly as many distinct ids**, and it has survived being restarted, having every backend restarted under it, and being moved from one system to another without losing a row. **Nothing was removed to reach these numbers**: the two required routes are additions, and registration, script login, sessions, end-to-end-encrypted rooms and the full `/api/*` surface all still work through the same URL and are still covered by the 69-assertion backend suite and the 39-assertion balancer suite, both of which pass. The previous assignment's URLs and direct ports keep working, its files on the lab systems are untouched, and the whole system can be redeployed, re-measured and demonstrated from the scripts in the repository.